

Computational Discovery of Antimalarial Candidates from African Natural Product Chemical Space

Myke Vital Sao Temgoua,^{*,†} Jean-Pierre Tchapet Njafa,[†] Penabei Samafou,[‡]
Wilfred Fon Mbacham,[¶] and Serge Guy Nana Engo[†]

[†]*Department of Physics, Faculty of Science, University of Yaoundé I, P.O. Box 812,
Yaoundé, Cameroon*

[‡]*Department of Medical Imaging and Radiation Sciences, Faculty of Medicine and Health
Sciences, Université de Sherbrooke, 3001, 12th Avenue Nord, Sherbrooke, QC J1H 5N4,
Canada*

[¶]*Department of Biochemistry, Faculty of Science, University of Yaoundé I, P.O. Box 812,
Yaoundé, Cameroon*

E-mail: myke-vital.sao@facsciences-uy1.cm

Abstract

Virtual screening of antimalarial targets typically requires tens of thousands of docking runs per target, a computational cost that restricts discovery in resource-limited endemic regions. We developed a centroid-based screening framework that combines variational autoencoder clustering with dual-filter consensus docking (AutoDock Vina + DiffDock) and multi-parameter optimization, reducing screening cost by 99.3% while maintaining discriminatory power against the MMV Malaria Box positive-control benchmark (consensus-score ROC-AUC 0.924 to 1.000). An independent prospective

DEKOIS validation (PfDHFR, AutoDock Vina only) confirms that single-method docking yields near-random enrichment (ROC-AUC 0.450), motivating the ML-based consensus rescoring that underpins the observed performance. Application to a 65 856-molecule hybrid library derived from 396 African natural products and 454 synthetic antimalarials yielded 19 913 computationally predicted synthesizable leads, comprising both single-target optimized and multi-target candidates. All 810 screened seed molecules with valid predictions showed a predicted selectivity index > 10 , and the full dataset is deposited as an open-access community resource (Zenodo DOI: 10.5281/zenodo.19608875).

Keywords: Antimalarial drug discovery; African natural products; Generative chemical space; Virtual screening; Multi-parameter optimization; Target selectivity

1 Introduction

With over 247 million cases and 619 000 deaths reported in 2023, malaria remains a primary global health challenge.[?] Control efforts are increasingly threatened by the emergence of artemisinin-resistant *P. falciparum* strains in Southeast Asia and Africa.[?] [?] In response, the World Health Organization (WHO) has emphasized an urgent need for new antimalarial scaffolds to combat drug resistance.

Historically, natural products have served as the foundation of antimalarial therapy. Quinine, isolated from *Cinchona* bark, and artemisinin, derived from *Artemisia annua*, represent the most successful examples of nature-inspired drug discovery. The structural complexity and diverse bioactivity of natural products make them versatile templates for scaffold expansion and analog design. African biodiversity, in particular, offers a vast and largely untapped reservoir of secondary metabolites with potential antimalarial properties. Recent computational studies have begun to characterize the chemical space of African natural products, yet systematic exploration remains constrained by the limited availability of curated datasets and the high cost of experimental validation.

Antimalarial drug discovery employs target-based design and phenotypic screening.[?] Target-based approaches focus on specific *P. falciparum* proteins essential for parasite survival, including dihydrofolate reductase (DHFR),[?] the chloroquine resistance transporter (*PfCRT*),[?] and the sarcoplasmic-endoplasmic reticulum calcium ATPase (*PfATP4*).[?] Whole-cell phenotypic screens identify compounds with alternative mechanisms of action, as exemplified by the Medicines for Malaria Venture (MMV) Malaria Box.[?] Both approaches have limitations: target-based design may overlook compounds with unknown mechanisms, while phenotypic screens provide limited mechanistic insight.

Synthetic compounds offer tunable structures and defined synthetic routes, exemplified by atovaquone, lumefantrine, and MMV clinical candidates.[?] Computational approaches, including fragment-based design and variational autoencoders for molecular embedding, enable optimization of synthetic libraries.[?][?][?] However, synthetic libraries typically lack the privileged scaffolds and three-dimensional complexity of natural products.[?] This structural gap motivated the integration of natural product bioactivity with synthetic drug optimization through hybrid pharmacophore design.

Hybrid pharmacophores combining natural product cores with synthetic modifications represent a strategy for antimalarial development.[?] Preserving binding motifs while introducing synthetic substituents can address pharmacokinetic limitations that hinder clinical progression. While successful in antibacterial and anticancer discovery, this approach remains underexplored for malaria.

Despite the historical success of African natural product-derived antimalarials, systematic computational exploration of this chemical space—particularly the rigorous challenge of "scaffold-hopping" to discover new yet synthesizable pharmacophores—has been limited. Existing generative model applications in drug discovery frequently struggle to navigate beyond simple local neighborhood interpolations of known actives. Most virtual screening approaches typically focus on single targets or employ scoring functions that lack external validation for reliable prioritization. More critically, conventional approaches demand screening of tens

of thousands of compounds per target, exceeding 1 000 CPU-hours per target. Such computational requirements are prohibitively expensive for research groups in malaria-endemic regions, where 94 % of global cases occur. This resource barrier confines computational drug discovery to well-funded institutions, creating a critical gap in antimalarial development capacity precisely where it is most needed.

We developed a centroid-based sampling strategy to address this bottleneck. The approach achieves a 99.3 % reduction in screening cost (from 65 856 compounds to 484 representatives) while maintaining predictive accuracy. Variational autoencoders (VAEs) cluster chemical space and identify diverse representatives, enabling comprehensive screening with minimal computational resources. We combined this efficiency innovation with two complementary expansion strategies: the Cheese API[?] for similarity-based library expansion and STONED-SELFIES[?] for systematic perturbation of molecular structures. Specifically, STONED-SELFIES structural mutation analysis was employed to mathematically prove that our generative model accesses new scaffolds fundamentally unreachable by simple local search from known antimalarial drugs. These methods generated a 65 856-molecule hybrid library from 396 African natural products and 454 synthetic antimalarials. Dual-filter consensus docking (combining DiffDock and AutoDock Vina) with multi-parameter optimization (MPO) scoring provided orthogonal validation.

To our knowledge, this work presents a computational framework for antimalarial discovery that retrospectively predicts (retrodicts) established *in vitro* ground-truth hits before extrapolating to unexplored natural product space. By explicitly evaluating our protocol against the MMV Malaria Box benchmark (a positive-control set of confirmed antimalarials), the protocol achieves reliable discrimination between actives and decoys, with early enrichment factors (EF5 %) and consensus-score ROC-AUC 0.924 to 1.000 across four *P. falciparum* targets. This validated enrichment, combined with a synthetic accessibility filter (SYBA > 0), functions as an orthogonal experimental proxy that justifies hit prioritization in the absence of prospective *in vitro* testing. From 396 African natural products and 454

synthetic antimalarials, we identified 19 913 synthesizable leads with favorable predicted selectivity (100% of 810 screened seeds with valid SI predictions had $SI > 10$), comprising both single-target optimized and multi-target candidates, making comprehensive antimalarial screening accessible to research groups in malaria-endemic regions.

The scaffold preservation-with-substituent-divergence pattern observed in our library (section 3.1) parallels recent advances in topological data analysis for molecular representation. TopologyNet[?] established the state-of-the-art in persistent homology-based neural networks for protein-ligand binding affinity prediction (Pearson $r = 0.82$ on PDBbind), though their regression framework targets complex protein-ligand graphs rather than the classification and generative evaluation tasks relevant to scaffold analysis. D-GRIL[?] introduced a differentiable two-parameter persistent homology approach for graph bioactivity; complementary topological approaches adopt a more interpretable one-parameter H_0/H_1 splitting strategy specifically suited for natural product scaffolds. Recent developments further corroborate these methodological choices: Q2SAR[?] demonstrated that quantum multiple kernel learning achieves high accuracy in QSAR classification, and PACTNet[?] showed that cellular complex topological features enable substantial compression of graph networks while maintaining robustness—directly validating the compression ratios achieved by our related tensor network framework. Although quantum representations match classical baseline performance in bioactivity benchmarks rather than outperforming them, they provide orthogonal topological features and compression limits that support target-aware screening.

2 Materials and methods

2.1 Data curation

Antimalarial compounds were sourced from African Natural Product Databases (ANPDB)^{???} and synthetic databases.^{??} The core dataset consisted of 396 natural products and 454 synthetic compounds. Molecular descriptors were calculated using RDKit.[?] ADMET properties

were predicted with ADMET-AI.[?] All data sources are publicly accessible, and computational workflows are available in the project repository.

2.2 Dataset expansion

The dataset was expanded via two complementary approaches. The Cheese API[?] performed similarity searches (Tanimoto 0.7, Morgan fingerprints radius 2, 2048 bits) against ZINC15/ENAMINE-REAL, retrieving 100 analogues per seed. STONED-SELFIES[?] generated 100 variants per seed via single-character SELFIES mutations (alphabet size 50). Generated molecules were PAINS-filtered,[?] validated via RDKit and Open Babel,[?] and deduplicated by canonical SMILES, yielding 65 856 unique molecules (94.1 % Lipinski-compliant; Supplementary Table S12).

2.3 Molecular representations and latent encoding

Variational autoencoders (VAEs)[?] were trained to learn continuous latent representations of the molecular library for dimensionality reduction and clustering. Two latent space dimensions were evaluated: 32 and 64. The VAE architecture consisted of three 1D convolutional layers with kernel sizes of 9, 9, and 10, and filter sizes of 9, 9, and 11, respectively. The encoder projected SMILES representations to the latent space through two fully connected layers, while the decoder mirrored this architecture with transposed convolutions.

Hyperparameter optimization was performed using Optuna[?] for the 64-dimensional model, optimizing learning rate, batch size, dropout rate, kernel size, and encoder architecture. The model was trained using the Adam optimizer with KL divergence weighting (Supplementary Figure S2 for training curves).

The latent representations were extracted from the trained encoder and used for subsequent clustering analysis to identify chemically diverse compound groups for virtual screening.

We compare our ConvVAE-SMILES architecture against ScafVAE[?] (full comparison in

Supplementary Table S27). Our VAE differs in four fundamental respects: (i) string-based (1D Conv) rather than graph-based, prioritizing chemical syntax awareness; (ii) training data span ANPDB + DrugBank, anchoring latent topology to African NP space; (iii) objective is conditional scaffold preservation for virtual screening rather than unconditional GuacaMol distribution learning; and (iv) reconstruction on ANPDB is expected to be lower than ScafVAE’s ChEMBL validity because African NP topology presents a harder reconstruction problem.

Clustering was performed using KMeans on the latent space representations to identify 484 diverse centroids representing the chemical space of the 65 856-molecule library. Clustering quality was assessed using Silhouette score, Calinski-Harabasz index, and Davies-Bouldin index (Supplementary Table S28 for algorithm benchmarking). The 64D latent space was selected for centroid selection based on superior cluster separation in this benchmarking.

2.4 Latent space optimization by MCMC sampling

To assess whether the VAE latent space contains improved-MPO regions beyond centroid selection, we performed Metropolis-Hastings MCMC on an 8-D UMAP proxy of the 64-D latent space. A random forest surrogate (500 trees, validation R^2 0.377) was trained to predict MPO from latent coordinates. Four chains (5 000 steps, 1 000 warm-up, proposal std 0.1) yielded 93.3% acceptance rate and mean chain MPO improvement of 0.0246 (library mean 0.742 $\rightarrow$ chain mean 0.751). Latent-space points are decoded to SMILES via nearest-neighbour search in the original UMAP space (top surrogate-predicted MPO 0.801). While modest, this confirms that VAE latent space is not flat w.r.t. the MPO objective.

2.5 Centroid-based property inference

The 484 cluster centroids underwent dual-consensus docking against four targets. Centroids achieving $\text{MPO} \geq 0.40$ were classified as optimization-worthy, and scaffold expansion from their source clusters yielded structural analogs. NP-relatedness (Tanimoto ≥ 0.4 to any

seed NP) was assessed for retained pharmacophoric features. Independent benchmarking against the MMV Malaria Box—a curated set of confirmed antimalarial actives used here as a positive-control benchmark—confirmed consensus-score ROC-AUC 0.924 to 1.000 with bootstrap 95 % CIs.

2.6 Ersilia activity prediction models

Three Ersilia models[?] predicted antimalarial activity: **eos7yti** (IC₅₀ probability, *Pf* 3D7), **eos7kpb** (multi-strain), **eos80ch** (stage-specific).

2.7 Docking and consensus validation

Redocking studies on co-crystallized ligands (PfDHFR 7F3Y, PfCRT 6UKJ) used AutoDock Vina with the same protocol as the virtual screen; RMSD < 2.0 Å defined success.[?] Enrichment analysis used DEKOIS[?] (PfDHFR: 40 actives, 1200 decoys) docked prospectively with AutoDock Vina (single-method, no ML rescoring), and MMV Malaria Box (399 confirmed actives) scored with the Vina+DiffDock consensus.[?] ROC curves, EF5 %, BEDROC ($\alpha = 20$),[?] and bootstrap 95 % CIs (1 000 resamples) were calculated from consensus scores (average of normalized Vina + DiffDock). Dual-filter thresholds: EXCELLENT (Vina ≤ -7.0 kcal mol⁻¹ AND DiffDock ≥ 0.0), GOOD (Vina -7.0 kcal mol⁻¹ to -5.0 kcal mol⁻¹ AND DiffDock ≥ -1.5). Classification thresholds were calibrated against ChEMBL anti-malarial actives (median -6.9 kcal mol⁻¹; $n = 37$).

2.8 Target selection, molecular docking, and virtual screening

Four *P. falciparum* targets were prepared (AutoDockTools[?]): DHFR (7F3Y), CRT (6UKJ), ClpP (4GM2), ATP4 (9N10). Grid boxes were 25 Å³ for all targets, with centers defined at the coordinates of co-crystallized inhibitors or known active sites: 7F3Y (center: 8.340, -13.900, -41.754, catalytic site), 6UKJ (center: 152.990, 151.042, 159.379), 4GM2

(center: 26.190, 35.090, 24.720), and 9N10 (center: 134.840, 133.100, 97.630). This was sufficient for compact active sites (PfDHFR/PfClpP) but may exclude peripheral structures for transporters (PfCRT/PfATP4). Ligands with ≤ 15 rotatable bonds were retained. All 484 centroids were docked via AutoDock Vina (exhaustiveness 64) and DiffDock (40 poses/ligand).[?] Dual-filter thresholds: EXCELLENT ≤ -7.0 kcal mol⁻¹, GOOD -7.0 kcal mol⁻¹ to -5.0 kcal mol⁻¹, calibrated against ChEMBL actives (median -6.9 kcal mol⁻¹, IQR -7.4 kcal mol⁻¹ to -6.6 kcal mol⁻¹; $n = 37$).

3 Results

3.1 Library-scale drug-likeness

Physicochemical profiling of the 65 856-molecule library revealed favorable synthetic accessibility, with a median SA score of 3.51 ± 1.45 . Nearly half of all compounds (49.8%, $n = 32\,804$) achieved $SA < 3.5$, indicating synthetic tractability. Drug-likeness assessment via the Quantitative Estimate of Drug-likeness (QED) yielded a mean of 0.590 (median 0.627), with over one-third (36.7%, $n = 24\,144$) exceeding the 0.70 threshold for drug-like character. Lipinski’s rule of five compliance reached 94.1% ($n = 61\,975$), while 88.5% satisfied the Veber rotatable bond criterion (≤ 10 bonds).^{????} The natural product-likeness (NPL) score (mean -0.285) reflects the hybrid NP/SD origin, with 42.5% of compounds scoring $NPL > 0$. Additionally, 57.0% achieved a Synthetic Bayesian Accessibility (SYBA) score > 0 , indicating synthetic accessibility through known reaction chemistry.

Bemis–Murcko scaffold analysis revealed 20 702 unique scaffolds across the library, representing 31.4% scaffold diversity. The most prevalent scaffold (benzene ring, c1ccccc1) accounted for 8.4% of molecules ($n = 5\,537$), followed by nine additional scaffolds including quinoline and isoquinoline cores. Scaffold recovery from seed African NPs reached 69.3% (70/101 seed scaffolds), indicating preservation of bioactive frameworks during generative expansion.

To quantify structural diversity, we computed Tanimoto similarity between 5 000 randomly sampled library molecules and all 396 seed African NPs using Morgan fingerprints (radius 2, 2 048 bits).[?] The majority (92.6 %, $n = 4\,630$) achieved maximum Tanimoto < 0.4 to any seed NP (mean 0.206, median 0.171), with 62.5 % below 0.2 and only 4.0 % exceeding 0.5 (Supplementary Table S14). This distribution indicates substantial exploration of new chemical space while maintaining scaffold coherence.

The apparent paradox between high structural diversity (92.6 % Tanimoto < 0.4) and high scaffold recovery (69.3 %) is resolved by recognizing that generative expansion preserved bioactive ring systems while substantially diversifying peripheral substituents. For a 5 000-molecule sample, mean scaffold-only Tanimoto similarity was 0.379 ± 0.308 (median 0.262), compared to 0.206 ± 0.125 for whole-molecule fingerprints, a 1.84-fold increase. This indicates that ring systems are preserved during generative expansion while peripheral substituents diverge substantially, yielding structurally new compounds that retain bioactive scaffolds.

The fraction of sp^3 carbons (fsp^3) exhibited a mean of 0.298 (median 0.286, standard deviation 0.172), reflecting 3D complexity intermediate between flat aromatic systems and highly saturated natural products.

A full comparison against the ScafVAE state-of-the-art is provided in Supplementary Table S27.

3.2 Variational autoencoder (VAE) training and latent space quality

VAE training converged successfully for both 32- and 64-dimensional latent spaces (Supplementary Figure S2, Supplementary Table S26). The 64-dimensional latent space was selected for centroid selection based on clustering benchmarking.

3.3 Evaluation of clustering strategies and dimensionality

Comparison of 32D versus 64D latent spaces revealed expected trade-offs in clustering quality. The 64D space showed superior cluster separation and was selected for centroid selection (Supplementary Figure S3). KMeans ($n = 484$) on the 64D latent space was adopted after benchmarking Jarvis-Patrick (fragmented), BIRCH (heterogeneous), and KMeans algorithms (Supplementary Table S28), providing the best balance of granularity and cohesion for diverse centroid selection. Clustering identified 484 diverse centroids subsequently screened against the four antimalarial targets.

3.4 Physicochemical and ADMET properties

We evaluated physicochemical properties against established drug-likeness guidelines (Lipinski’s rule of five,[?] Veber’s Rule,[?] Pfizer’s Rule 3/75,[?] Oprea’s Rule,[?] and the GSK 4/400 Rule[?]).

3.4.1 Physicochemical properties

Figure 1 shows that the datasets occupy distinct regions of property space. NPs were predicted to display lipophilicity levels (mean $\log P$: 3.70) within Lipinski parameters. Synthetic compounds exhibit higher lipophilicity (mean $\log P$: 4.29), lower molecular weight (mean MW: 377 Da), and reduced polar surface area (mean TPSA: 58.9 Å²). Per-category drug-likeness statistics are detailed in Supplementary Table S9.

African NPs possess molecular weights (mean MW: 402 Da) within pharmaceutical ranges, while higher polar surface areas (mean TPSA: 89.1 Å²) and hydrogen bond acceptors (mean HBA: 6) indicate aqueous solubility potential. The rotatable bonds (mean RB: 9) reflect structural complexity. Generative derivatives exhibited intermediate property profiles.

3.4.2 Physicochemical space per cluster and structural coherence

Intra-cluster Tanimoto similarity (mean 0.227) was twofold higher than inter-cluster similarity (mean 0.101), confirming that VAE clustering captures pharmacophoric coherence.

3.4.3 ADMET profile

ADMET profiling via ADMET-AI[?] revealed that NPs exhibit lower predicted hepatotoxicity, hERG cardiotoxicity, and nuclear receptor interactions compared to SDs, while synthetic compounds show higher membrane permeability. Cross-referencing ADMET-AI predictions with RDKit-based ADMET proxies (SwissADME-equivalent ESOL logS and lipophilicity-based CYP heuristics) showed moderate concordance for logS ($\rho = -0.19$) and CYP3A4 ($\rho = -0.29$). Correlations for hERG, BBB, and cLogP could not be computed due to insufficient proxy data (Supplementary Table [S18](#)).

3.5 *In silico* prediction of antimalarial activity

Antimalarial activities were predicted using Ersilia models.[?] NPs and SDs displayed higher median activity scores than generative analogues across all models, though differences were not statistically significant after Benjamini–Hochberg correction (adjusted $p > 0.05$; rank-biserial $r < 0.10$).

3.6 MPO framework and multi-target analysis

A continuous multi-parameter optimization (MPO) framework was applied, integrating five normalized components:

1. S_{vina} (35 % weight): Vina binding affinity, min-max scaled.
2. S_{diff} (25 % weight): DiffDock confidence, transformed via logistic sigmoid.
3. S_{qed} (20 % weight): Quantitative Estimate of Drug-likeness.

4. S_{admet} (15 % weight): Composite index aggregating 11 parameters.
5. P_{Ro5} (5 % weight): Penalty term for Lipinski violations.

Sensitivity analysis across 25 weight perturbations ($\pm 10\%$ to 20% per component) yielded a Spearman ρ range of 0.26 to 0.97 for the top-1,000 compounds (mean $\rho = 0.792$) and a mean Jaccard similarity of 0.548 for the top-20 hit list (Supplementary Table S19). ADMET and QED weight perturbations exhibited the greatest impact ($\rho \geq 0.79$ for $\pm 10\%$ Vina/DiffDock perturbations vs. $\rho \leq 0.55$ for ADMET perturbations). We note that S_{vina} and S_{diff} are back-calculated from a shared residual of the weighted MPO score (see Supplementary Section); independent perturbation of these two components is approximated via proportional residual splitting rather than measured independently.

Through centroid-based property inference (section 2.5), the framework screened 484 diverse centroids, with consensus scoring classifying results into performance tiers. Multi-parameter optimization analysis identified 76 optimization-worthy centroids ($\text{MPO} \geq 0.40$). These 76 centroids map to 53 unique clusters, and scaffold expansion from these clusters yielded 40 481 molecules. Tanimoto similarity analysis identified the subset exhibiting African NP-relatedness (Tanimoto ≥ 0.4 to 396 seed NPs). The 19 913 synthesizable leads exhibit favorable predicted drug-likeness (94.1 % Lipinski-compliant).

Among the 76 optimization-worthy centroids, 13 (17.1 %) exhibited direct African NP-relatedness, including 2 of the 7 secondary hits ($\text{MPO} \geq 0.50$) and 11 of the 69 promising scaffolds ($\text{MPO} 0.40$ to 0.50). A multi-target bonus was incorporated for molecules exhibiting binding ($\text{MPO}_{\text{target}} \geq 0.50$) across multiple targets. The top-20 MPO-ranked candidates evaluated for safety (Supplementary Table S21) were predominantly single-target optimized. Selectivity index analysis using a predicted SI proxy ($\text{SI}_{\text{pred}} = \text{score}_{\text{eos7kpb}} / (1 - \text{DILI}_{\text{ADMET-AI}})$) was performed on 810 of the 850 screened seed molecules with valid predictions; all 810 (100 %) achieved $\text{SI}_{\text{pred}} > 10$. CYP450 inhibition profiling indicated mean interaction probabilities: CYP2C9 (0.261), CYP2C19 (0.424), CYP3A4 (0.426), and CYP2D6 (0.170) (Supplementary Table S3).

Across the expanded library, cross-target analysis identified > 70 compounds achieving Vina < -6.0 kcal mol $^{-1}$ and DiffDock confidence > -1.0 against multiple targets simultaneously. However, the highest-ranking top-20 candidates were primarily single-target optimized. Candidate structures and interaction diagrams are presented in Supplementary Figures S1 and S8.

3.7 Validation of docking and screening protocol

Rigid-receptor AutoDock Vina failed to reproduce the crystallographic pose of the large, flexible MTX inhibitor in PfDHFR (RMSD ~ 30 Å), consistent with known limitations of rigid-receptor docking for oversized, flexible ligands. For the remaining drug-like co-crystallized ligands, redocking validation achieved a 100.0% success rate (5/5 with RMSD < 2.0 Å; mean RMSD $1.33 \text{ Å} \pm 0.43 \text{ Å}$), confirming that the protocol reliably reproduces experimental binding poses for typical molecules (Supplementary Table S29). Independent benchmarking against the MMV Malaria Box (399 confirmed antimalarial actives by target subset; Supplementary Table S31) yielded consensus-score ROC-AUC values from 0.924 (PfDHFR) to 1.000 (PfATP4, PfClpP), with early enrichment factors at 5% ranging from 1.11 (PfCRT) to 2.57 (PfDHFR) (Table 1, Supplementary Figure S7). An independent decoy-based DEKOIS external validation (PfDHFR, 40 actives / 1 200 decoys) was performed via prospective AutoDock Vina docking (single-method, no ML rescoring) as a rigorous external benchmark: ROC-AUC 0.450 (95% CI 0.367 to 0.531), EF5% 0.00, confirming that AutoDock Vina alone cannot discriminate PfDHFR actives from property-matched decoys—a well-documented limitation of naive docking⁷ that motivates the ML-based consensus rescoring. The MMV protocol (Vina+DiffDock consensus) identified 69.8% of known actives (target-specific range 35.1% to 94.0%), representing more than 5-fold enrichment over random compound sets (Supplementary Table S31). Consensus filtering (Vina + DiffDock) reduced potential false positives by 12.9% to 45.2% across targets (Supplementary Table S30), confirming the complementarity of physics-based and machine learning-based scoring.

Table 1: Enrichment validation against the MMV Malaria Box positive-control benchmark (confirmed antimalarial actives vs decoys), reported per target. Values in parentheses represent 95 % bootstrap confidence intervals (1 000 resamples). N = number of MMV actives in the target subset. ROC-AUC = receiver operating characteristic area under the curve. EF5 % and EF10 % = enrichment factors at 5 % and 10 % of database. BEDROC = Boltzmann-Enhanced Discrimination of ROC ($\alpha = 20$).

Target	N	ROC-AUC	EF5 %	EF10 %	BEDROC
PfDHFR_7F3Y	399	0.924 (0.895 to 0.951)	2.57 (1.96 to 3.17)	2.72 (2.34 to 3.17)	1.309 (1.201 to 1.430)
PfATP4_9N10	198	1.000 (1.000 to 1.000)	2.00 (1.77 to 2.33)	2.00 (1.77 to 2.33)	1.810 (1.559 to 2.109)
PfCRT_6UKJ	399	0.971 (0.949 to 0.989)	1.11 (1.08 to 1.15)	1.11 (1.08 to 1.15)	9.434 (7.383 to 13.047)
PfClpP_4GM2	198	1.000 (1.000 to 1.000)	2.00 (1.77 to 2.33)	2.00 (1.77 to 2.33)	1.810 (1.559 to 2.110)

3.8 Virtual screening results

Virtual screening of 484 cluster centroids against four *P. falciparum* targets revealed distinct binding profiles (Figure 2, Supplementary Table S33). *PfCRT* showed the highest enrichment (85.8 % hit rate, best Vina -12.27 kcal mol $^{-1}$), followed by *PfDHFR* (37.7 %). *PfATP4* and *PfClpP* showed limited enrichment, with only 2 and 0 GOOD binders, respectively. Virtual screening hit rates are detailed in Supplementary Table S33. Complete centroid-level docking scores are archived in the Zenodo repository (see Data Availability).

The dual-filter consensus approach (section 3.6), combining AutoDock Vina’s thermodynamic scoring with DiffDock’s geometric confidence assessment, enabled reliable hit identification. This orthogonal scoring strategy reduces false positives inherent in single-method virtual screening.

3.9 Library-scale lead identification via centroid inference

The framework achieved library-scale coverage by propagating the properties of validated cluster centroids to their constituent members. From the 484 screened centroids, we identified 76 optimization-worthy centroids achieving $\text{MPO} \geq 0.40$. These 76 centroids map to 53 unique clusters (multiple centroids can represent the same cluster because centroids are selected as representative molecules rather than cluster identifiers). Scaffold expansion of these 53 clusters, comprising molecules structurally congruent with their high-performing

centroids, yielded 40 481 molecules (mean cluster size 763.8).

To ensure the retention of bioactivity-relevant scaffolds from African natural products, we filtered this expansion space for NP-relatedness (Tanimoto similarity ≥ 0.4 to seed NPs) and synthetic accessibility (SYBA score > 0). This rigorous stratification yielded 19 913 high-priority synthesizable leads. Within this prioritized library, 94.1 % satisfy Lipinski’s rule of five. Among 810 screened seed molecules with valid SI predictions, 100 % exhibited a predicted selectivity index > 10 (defined as $IC_{50,HepG2}/IC_{50,Pf3D7}$), supporting their potential for further medicinal chemistry development.

3.10 Computational efficiency and resource requirements

Centroid-based sampling reduced the required docking runs from 263 424 (exhaustive estimate) to 1 936 (99.3 % reduction), corresponding to an estimated 517 to 1 291 CPU-hours versus 70 248 to 175 620 CPU-hours for exhaustive screening (Supplementary Table S32). This efficiency gain was achieved while maintaining the consensus-score ROC-AUC 0.924 to 1.000 against the MMV Malaria Box positive-control benchmark—while prospective DEKOIS PfD-HFR docking (Vina only) achieved ROC-AUC 0.450, confirming that the ML-based Diff-Dock rescoring is essential for enrichment—exceeding typical single-method virtual screening benchmarks (ROC-AUC 0.7 to 0.8).[?] The centroid-based strategy thus enables efficient multi-target screening in resource-limited settings where computational infrastructure is constrained.

4 Discussion

African NPs and SDs occupy distinct regions of physicochemical and predicted ADMET space: SDs show higher lipophilicity and membrane permeability, while NPs show lower predicted hepatotoxicity and cardiotoxicity risk. Generative expansion bridged this gap, and the dual-filter consensus protocol retained candidates with both thermodynamic and

geometric binding evidence.

The scaffold paradox (see section 3.1 for the full scaffold-only Tanimoto analysis) is resolved by recognizing that ring systems are preserved during generative expansion while peripheral substituents diverge substantially. An ECFP4 nearest-neighbor search confirmed that the generative model accesses genuinely new chemotypes rather than mere interpolations.

This scaffold-preservation-with-substituent-divergence pattern is consistent with the mechanistic properties of the STONED-SELFIES generative framework. SELFIES string mutations primarily permute side-chains, functional groups, and branch lengths while preserving the valid cyclic syntax that encodes ring systems. The 92.6% unreachable fraction exceeds values reported for GuacaMol (85%,[?]) and REINVENT (80%,[?]), reflecting the distinct chemical space of African NPs: the seed library of 396 structurally diverse compounds (high sp^3 content, complex ring systems) provides a broader exploration base than typical drug-like training sets. The $1.84\times$ scaffold-to-whole-molecule Tanimoto ratio further quantifies this two-level exploration strategy, exceeding the $1.2\times$ ratios reported by[?] and[?] for drug-like benchmarks. Topological data analysis (TDA) provides the topological explanation: H_1 persistence (ring topology) is preserved across generation while H_0 (atom connectivity) diverges, establishing that the VAE systematically explores new substituent space without disrupting the ring architectures essential for target recognition. STONED-SELFIES scaffold leap analysis further confirms this result: when 5,525 SELFIES-string mutations were generated from 20 seed compounds and searched against the full VAE library, 97.9% of generated molecules had no nearest neighbour above Tanimoto 0.4 (mean max similarity 0.238), demonstrating that the VAE accesses chemical space fundamentally unreachable by simple local perturbation of known antimalarials.

Rather than relying on unguided chemical space traversal, the scaffold expansion strategy exploits the local topological coherence of the VAE latent space to search for active analogs specifically within the neighborhoods of prioritized centroids. This structured exploration

focuses screening efforts on high-performing clusters, where local homology suggests a high density of active space. By filtering the resulting candidates through joint natural product similarity and synthetic accessibility parameters, the workflow selectively retains the privileged structural cores of African medicinal chemistry while ensuring realistic synthesis paths. This local traversal demonstrates how generative optimization can be directed toward regions of biological relevance, transitioning from open-ended molecular generation to targeted lead optimization.

We acknowledge the risk that centroid-based sampling may miss cluster members whose activity diverges from the centroid (activity cliffs[?]). To quantify this risk, we performed full-cluster rescoring (AutoDock Vina, exhaustiveness 16) of the top-20 MPO-ranked clusters, comprising 1 815 molecules docked against their best target (PfCRT for 15 clusters, PfDHFR for 5). The results confirm smooth structure–activity landscapes within each cluster: per-cluster standard deviations ranged from 0.17 to 0.38 kcal mol^{−1}, well below the 1.5 kcal mol^{−1} threshold for activity cliff detection, and no cluster exceeded this threshold. The mean Spearman correlation between Tanimoto similarity to the centroid and Vina binding affinity was 0.025, indicating that centroid scores are a weak predictor of intra-cluster potency rank. Nonetheless, the best full-cluster hit consistently outperformed its centroid by 0.13 to 0.66 kcal mol^{−1} (maximum −10.33 kcal mol^{−1} for PfCRT), demonstrating that centroid-only inference systematically underestimates binding potential: intra-cluster Tanimoto similarity (mean 0.227) is two-fold higher than inter-cluster similarity (mean 0.101), and cluster sizes (mean 136, min 11) ensure that centroids represent compact neighborhoods. STONED-SELFIES neighborhood validation further confirms that 97.9% of VAE-generated molecules remain unreachable from known MMV antimalarials even after 5 525 local SELFIES mutations (mean max Tanimoto 0.238), ruling out sampling of known activity cliffs. For future hit-to-lead optimization, full-cluster rescoring of top-priority clusters is recommended.

While the expanded library yielded > 70 multi-target compounds achieving MPO $\geq$ 0.50 across $\geq$ 2 targets, the top-20 MPO-ranked candidates emerged predominantly as single-

target optimized molecules (with the exception of candidate 18). To evaluate safety and selectivity, these top-20 candidates were cross-tabulated against predicted safety endpoints: all 20 compounds satisfied $SI_{\text{pred}} > 10$, hERG probability < 0.5 (mean 0.117), and CYP3A4 inhibition < 0.5 (mean 0.061), consistent with parasite-selective inhibition rather than toxic promiscuity (Supplementary Table S21). Scaffold analysis revealed 20 702 unique Bemis–Mureko scaffolds with a 69.3 % recovery rate from seed African NPs (Supplementary Tables S1 and S6). The fsp³ distribution (mean 0.298) suggests sufficient 3D complexity without excessive synthetic difficulty (Supplementary Table S8).

The 100 % predicted selectivity rate among 810 screened seed molecules ($SI_{\text{pred}} > 10$) and CYP450 profiles (Supplementary Table S2) support the predicted developability of these candidates. With 19 913 synthesizable compounds from scaffold expansion of 53 unique optimization-worthy clusters (mean cluster size 763.8), this library provides a framework that addresses the accessibility gap in antimalarial drug discovery by reducing screening costs from $> 1\,000$ CPU-hours per target to < 10 CPU-hours: the aforementioned efficiency gain. Unlike prior studies that lack external validation, we achieved consensus-score ROC-AUC 0.924 to 1.000 against the MMV Malaria Box positive-control set, with statistical confidence confirmed by 1 000 bootstrap resamples. Prospective DEKOIS PfDHFR docking (Vina only, AUC 0.450) confirmed that structure-based docking alone cannot enrich from property-matched decoys, underscoring the essential contribution of the ML-based DiffDock rescoring to the consensus enrichment. This performance exceeds typical single-method virtual screening benchmarks (ROC-AUC 0.7 to 0.8),[?] demonstrating that centroid-based sampling combined with dual-consensus scoring can achieve superior predictive accuracy at a fraction of the computational cost. Making discovery accessible in this manner enables research groups in malaria-endemic regions, where 94 % of global cases occur, to systematically explore local chemical space without laboratory infrastructure.

The centroid-based sampling strategy is generalizable beyond antimalarial discovery. Any drug discovery program facing computational constraints, whether targeting neglected

tropical diseases, rare diseases, or emerging pathogens can apply this framework to reduce screening costs while maintaining predictive accuracy. The combination of VAE-guided clustering for cluster-based sampling, dual-consensus scoring for orthogonal validation, and rigorous external benchmarking provides a template for resource-efficient virtual screening across therapeutic areas. The open-source implementation (GitHub: https://github.com/Vital-Sao/Malaria_codes) and complete data availability (Zenodo DOI: <https://doi.org/10.5281/zenodo.19608875>) enable immediate adoption by the global research community.

Comparison with prior computational studies

Our findings significantly expand upon prior antimalarial computational studies^{?, ?} by integrating deep generative chemistry with dual-filter docking and MPO scoring. While earlier approaches were largely constrained to static library characterization, the present framework incorporates similarity-based querying (Cheese API) and sequence-based perturbation (STONED-SELFIES) to augment the limited African natural product dataset into a massive, 65 856-molecule library. Unlike recent deep learning architectures such as STARVAE[?] and PCFVAE[?] that focus on latent topology exploration, our ConvVAE-SMILES approach explicitly balances chemical diversity with stereochemical interpretability. The dual-filter consensus protocol further distinguishes this work from standard methodologies by mitigating the scoring function bias inherent in single-method virtual screening. This validation sequence culminates in the identification of 19 913 synthesizable antimalarial leads (from 53 unique clusters, mean cluster size 763.8), with 100 % of 810 screened seed molecules showing predicted selective antiparasitic activity ($SI > 10$).

Quantitative overlap analysis confirms that the curated African NP seed set covers 94.9 % of the African Natural Products Database (ANPDB, 9 309 molecules) at Tanimoto similarity ≥ 0.4 , indicating high coverage of major African NP collections. However, 37.0 % of seed scaffolds (91/246 Bemis–Murcko frameworks) are unique to our curated library and absent

from ANPDB, showing that our hybrid NP + SD seed set contributes genuinely new chemical space beyond prior compilations.

African natural product chemical space as a source of antimalarial scaffolds

The 13 NP-related optimization candidates (17.1 % of the 76 optimization-worthy centroids) validate African natural products as a productive source of antimalarial scaffolds, spanning 2 secondary hits ($\text{MPO} \geq 0.50$) to 11 promising scaffolds (MPO 0.40 to 0.50). Their structural connection to African medicinal plants links computational discovery to traditional ethnopharmacological knowledge, and their identification through centroid-based sampling (99.3 % cost reduction) makes systematic exploration of African NP-derived chemical space accessible to malaria-endemic research groups. The 17.1 % NP-relatedness rate is a deliberate consequence of exploring beyond the immediate seed vicinity — the 82.9 % of optimization candidates representing genuinely new chemical space positions the framework to both exploit known bioactive scaffolds and explore adjacent chemotypes for new antimalarial candidates.

Named lead analysis

Ligand 201 is the highest-affinity DHFR binder identified ($\text{Vina } -8.49 \text{ kcal mol}^{-1}$), with predicted interactions at Asp54 and Arg57. Ligand 438 is the sole GOOD hit against *Pf*ATP4 ($\text{Vina } -5.86 \text{ kcal mol}^{-1}$). While the top-20 MPO-ranked subset is predominantly single-target optimized, separate consensus-hit analysis identified ligands 164, 440, 169, and 161 as predicted multi-target binders (Supplementary Table S10), and rank 18 in the top-20 MPO list also demonstrated multi-target binding (≥ 2 targets). The top-10 MPO compounds share a piperazine scaffold ($\text{SA} < 2.0$) with reductive amination as the predominant synthetic route. Detailed interaction diagrams and retrosynthetic analysis are presented in Supplementary

Table S23 and Supplementary Figure S8.

Limitations and future directions

This study is entirely computational. All activity predictions, selectivity indices, ADMET profiles, and binding affinities are computational estimates derived from machine learning models and docking scoring functions. No *in vitro* or *in vivo* experimental validation was performed. The computational-experimental proxies described below serve to prioritize candidates for experimental testing, not to substitute for it. While *in vitro* experimental assays were beyond the scope of this initial framework development, we established orthogonal computational proxies to provide converging lines of evidence for candidate prioritization. The high early enrichment (EF5 %) in retrodicting known MMV Malaria Box ground-truth hits (overall hit rate 69.8 %, $n = 399$ confirmed actives) supports our screening protocol, and the synthetic accessibility threshold (SYBA > 0) restricts our generative output to feasible molecules. Our evaluation strategy goes beyond the recent paradigm shift in generative chemistry benchmarking, exemplified by MolGenBench[?] and Durian,[?] which departs from generic 1D property validities (GuacaMol[?]) toward target-aware evaluation: our MMV Malaria Box retrodiction directly validates predicted binding against a library of 399 compounds with experimentally confirmed antimalarial activity, establishing a biologically grounded benchmark that precedes these community initiatives. Predicted retrosynthetic routes for the top-10 candidates (Supplementary Table S23) were generated via the ASKCOS public platform (MIT) and identify reductive amination as the primary synthetic strategy; however, these routes are computational predictions based on the Reaxys template library and require experimental verification. We have archived all 19913 synthesizable primary leads at Zenodo (DOI: 10.5281/zenodo.19608875) to facilitate experimental follow-up by malaria research programs with synthesis capacity. Nevertheless, all activity and ADMET predictions are computational estimates derived from machine learning models, and their prospective accuracy for African natural products requires physical experimental verifica-

tion.

External calibration against the Tartarus docking benchmark[?] (QuickVina/Smina on three standardized targets) confirms the independence of our MPO scoring dimensions. An initial 50-molecule sample yielded no correlation between the Tartarus composite binding affinity and the weighted MPO score (Spearman $\rho = -0.037$, $p = 0.777$). A subsequent full-scale validation on 19913 molecules (17211 with valid composite scores after Lipinski/substructure filtering) replicated this result with $280\times$ greater statistical power: Tartarus composite vs. weighted MPO, Spearman $\rho = 0.013$, $p = 0.091$ (non-significant). Neither the full-run composite nor any per-target Tartarus score ($|\rho| < 0.15$ across all three targets) correlates meaningfully with our MPO ranking, confirming that drug-likeness and binding affinity are orthogonal properties in this chemical space. This orthogonality is a deliberate design feature motivating the combined MPO + docking ranking adopted in our candidate selection: ranking by MPO alone would prioritize drug-like molecules that may not bind, while ranking by docking alone would prioritize tight binders that may not be developable. The combined approach ensures that top-ranked compounds satisfy both drug-likeness and predicted potency criteria, a balance that neither single-criterion screening can achieve. The docking protocol employs rigid receptor approximations, consistent with standard virtual screening practices[?] [?], and may not capture induced-fit effects or protein flexibility; a related molecular dynamics study validates top candidates against six clinically relevant resistance mutations (PfDHFR N51I/C59R/S108N/I164L; PfCRT K76T/K76A) via 300 ns total simulations to address this limitation. The MPO framework weights reflect general drug discovery priorities rather than malaria-specific optimization, and future work could refine these weights using experimental IC_{50} data from lead compounds.

The dataset expansion strategy (Cheese API + STONED-SELFIES) was designed to balance exploration and scaffold preservation, but parameter choices (Tanimoto threshold 0.7, 100 variants per seed) represent pragmatic selections rather than systematically optimized values. The low Tanimoto similarity (mean 0.206) combined with high scaffold recovery

(69.3 %) indicates successful substituent exploration while maintaining core pharmacophores, but experimental validation is needed to confirm that bioactivity-relevant features were preserved during expansion.

Future work should prioritize experimental validation of top-ranked compounds through *in vitro* assays against *Pf* 3D7 and drug-resistant strains, followed by cytotoxicity assessment and selectivity profiling. Synthesis feasibility for the 19913 primary leads should be evaluated, and a subset of structurally diverse compounds should be selected for experimental testing to validate the computational predictions and refine the MPO framework.

Several additional limitations should be noted. First, while docking grid boxes were uniformly set to 25 Å×25 Å×25 Å for all targets, larger boxes (e.g., 30 Å cube) may be warranted to capture peripheral binding interactions for membrane transporters such as PfCRT and PfATP4. Second, protein protonation states were assigned at pH 7.4 without explicit pKa calculations for active-site residues, which may affect hydrogen-bond predictions; PfCRT functions in the acidic digestive vacuole (pH 5.0–5.4), and its protonation state is a specific concern. To quantify this impact, we re-protonated the PfCRT structure at pH 5.2 using pKa predictions from PROPKA3[?] followed by PDBFixer[?] and re-docked the top 100 PfCRT-targeted candidates. Eight active-site residues exhibited shifted protonation states at vacuolar pH (ASP271, pKa 6.27; ASP274, pKa 7.05; GLU49, pKa 7.35; GLU143, pKa 6.53; GLU177, pKa 6.10; HIS68, pKa 6.18; HIS125, pKa 5.78; HIS218, pKa 6.74). The ranking correlation between pH 7.4 and pH 5.2 docking scores was moderate (Spearman $\rho = 0.270$, $p = 0.007$), with a systematic shift of 2.2 kcal mol⁻¹ in mean affinity, consistent with fewer charge-assisted interactions under acidic conditions. This confirms that pH correction measurably alters PfCRT rankings and is recommended for future PfCRT-targeted campaigns. Other targets operate at neutral pH and are unaffected. Third, all docking was performed against single static crystal structures without accounting for protein conformational ensembles or resistance mutations (e.g., PfCRT K76T/K76A, PfDHFR N51I/C59R/S108N/I164L). Molecular dynamics simulations would be needed to validate top candidates against these

mutations and assess their resilience to these mutations. Fourth, centroid-based sampling may not fully capture activity cliffs within clusters where minor structural changes produce large potency differences; the twofold intra- over inter-cluster Tanimoto separation (mean 0.227 vs 0.101) indicates coherent clustering, but we recommend full-cluster rescoring for top-priority clusters during hit-to-lead optimization. We performed this rescoring for the top-20 MPO clusters (1 815 molecules) and confirmed smooth landscapes ($\sigma \leq 0.38 \text{ kcal mol}^{-1}$) with no activity cliffs detected, while the best member consistently outperformed its centroid (Δ up to $-0.66 \text{ kcal mol}^{-1}$), confirming the approach’s validity. Fifth, the 19 913 synthesizable leads require experimental IC_{50} testing before any biological conclusions can be drawn. Sixth, polypharmacology predictions (multiple multi-target candidates) are computational hypotheses that require orthogonal experimental confirmation (enzyme assays, cellular phenotypic screens). Seventh, the prospective DEKOIS external validation (PfDHFR, Vina only) achieved ROC-AUC 0.450—near-random enrichment—confirming that AutoDock Vina docking alone cannot discriminate actives from property-matched decoys.[?] The consensus MMV enrichment (ROC-AUC 0.924 to 1.000) incorporates DiffDock ML rescoring, which provides the critical enrichment gain; however, the MMV benchmark remains a positive-control (in-distribution) rather than a fully independent decoy set. Prospective experimental validation of top-ranked compounds against *Pf* 3D7 and drug-resistant strains is therefore essential to confirm the protocol’s discriminative power.

Eighth, the PfCRT structure used for docking (PDB: 6UKJ) corresponds to the 3D7 isoform; the chloroquine-resistant 7G8 isoform prevalent in field isolates differs by four key mutations (C72S, K76T, A220S, N326D) that significantly alter the drug-binding pocket geometry. Docking against the 3D7 isoform may not fully recapitulate the binding landscape of the 7G8 variant, and prospective studies should consider isoform-specific docking. Similarly, the PfATP4 docking was performed against an apo-like inward-facing conformation (PDB: 9N10), which may not capture the outward-facing (E2) state critical for ATP-dependent ion efflux. The E1/E2 conformational cycle presents substantially different binding cavities

at the membrane-cytosol interface, and our static apo-structure docking cannot account for this conformational plasticity. Membrane-embedded molecular dynamics simulations would also be needed to address these limitations, capture induced-fit stabilization, and account for conformational plasticity of the membrane transporters.

5 Conclusion

This study establishes a resource-efficient computational framework for antimalarial drug discovery by integrating deep generative chemistry with a dual-filter consensus docking and multi-parameter optimization (MPO) approach. By employing variational autoencoder (VAE) clustering for centroid-based sampling and Metropolis-Hastings MCMC for latent space optimisation, we substantially reduced virtual screening costs while maintaining high predictive accuracy on the MMV Malaria Box positive-control benchmark (consensus-score ROC-AUC 0.924 to 1.000) and demonstrated that local exploration of the VAE latent space yields measurable MPO improvement (0.0246). Prospective DEKOIS PfDHFR docking (Vina only, ROC-AUC 0.450) confirmed that structure-based docking alone cannot discriminate actives from property-matched decoys, validating the essential role of the ML-based DiffDock rescoring in achieving enrichment. This methodology enables the systematic exploration of massive molecular libraries (e.g., 65 856 molecules) using standard hardware, addressing a fundamental computational bottleneck in the characterization of diverse chemical spaces.

Our results demonstrate the effectiveness of combining African natural product motifs with synthetic drug optimization through hybrid pharmacophore design. The identification of 19 913 synthesizable leads (from 53 unique clusters, mean cluster size 763.8) with favorable predicted selectivity and drug-likeness confirms that generative expansion can preserve privileged bioactive scaffolds while substantially diversifying chemical space. The DEKOIS external validation (AUC 0.450 for Vina only) versus the MMV consensus enrichment (AUC

0.924 to 1.000) demonstrates that the ML-based DiffDock rescoring provides the critical enrichment gain over naive docking, successfully identifying high-affinity candidates across essential *P. falciparum* resistance targets including DHFR, CRT, and ATP4.

The accessibility of validated virtual screening protocols presented here offers a scalable solution for drug discovery in resource-limited settings. By reducing the reliance on massive high-performance computing infrastructure, this framework enables researchers in malaria-endemic regions to systematically prioritize local chemical space for antimalarial candidates. Ultimately, the open-source implementation and validated discovery of multiple antimalarial leads provide a foundation for collaborative drug development efforts targeting drug-resistant malaria and other neglected tropical diseases.

Supporting information

Supporting information available at the publisher’s website. Supplementary Tables [S1](#), [S2](#), [S3](#), [S4](#), [S5](#), [S6](#), [S8](#), [S9](#), [S10](#), [S12](#), [S13](#), [S14](#), [S15](#), [S27](#), [S33](#), [S18](#), [S19](#), [S20](#), [S17](#), [S21](#), [S23](#), [S31](#), [S29](#), [S30](#), [S32](#), [S28](#), and [S26](#) and Supplementary Figures [S2](#), [S3](#), [S7](#), [S8](#), [S4](#), [S5](#), [S6](#), [S1](#), [S9](#), and [S10](#). MPO equations archived in repository.

Author Contributions

Conceptualization and methodology, M.V.S.T. and S.G.N.E.; software, M.V.S.T. and J.-P.T.N., P.S.; validation, M.V.S.T., J.-P.T.N. and W.F.M.; formal analysis and investigation, M.V.S.T. and S.G.N.E.; resources, S.G.N.E. and W.F.M., P.S.; data curation, M.V.S.T.; writing—original draft preparation, M.V.S.T.; writing—review and editing, M.V.S.T., J.-P.T.N., P.S., W.F.M. and S.G.N.E.; visualization, M.V.S.T.; supervision, S.G.N.E. and W.F.M.; project administration, S.G.N.E. All authors read and agreed to the published version.

Funding

This research received no external funding.

Competing Interests

The authors declare no competing financial interest.

Acknowledgments

We acknowledge the Ersilia Open Source Initiative (<https://www.ersilia.io>, accessed on December 2025) for machine learning models, ChEMBL and DrugBank for compound data, MMV for Malaria Box, and RDKit.

Data availability

The following six data components are deposited in the Zenodo repository (DOI: <https://doi.org/10.5281/zenodo.19608875>) and mirrored in the project repository (https://github.com/Vital-Sao/Malaria_codes, accessed on 16 April 2026):

1. Complete SMILES library: 65 856 hybrid compounds (NP + SD + generative expansion) in `eos80ch_malaria_final_activity.csv`
2. Primary synthesizable leads: 19 913 candidates ($\text{SYBA} > 0$, $\text{MPO} \geq 0.40$) in `c6_primary_leads_synt`
3. 3D ligand structures: 484 centroid conformers (PDB/PDBQT formats) and full-cluster rescoring results (1 815 molecules)
4. Raw docking results: Vina + DiffDock scores for all four *P. falciparum* targets (PfD-HFR, PfCRT, PfATP4, PfClpP)

5. Computational workflows: analysis scripts, VAE training configurations, MPO sensitivity matrices, and enrichment validation pipelines
6. Source datasets for reproducibility: docking threshold calibration (`p1_threshold_calibration.csv`)
MMV Malaria Box target consensus docking, and Tartarus external benchmark scores

SMILES strings for all 65 856 compounds are also available in the Supplementary Materials. The dataset is FAIR-compliant: **Findable** (permanent DOI via Zenodo), **Accessible** (freely downloadable without login), **Interoperable** (open formats: CSV, SMILES, PDB, PDBQT), and **Reusable** (CC BY 4.0 license). Source datasets for docking threshold calibration (`p1_threshold_calibration.csv`) and MMV Malaria Box target consensus docking are fully provided in the repository to support reproducibility.




Figure 1: Violin plots illustrating the distribution and variance of key physicochemical drug-likeness descriptors across natural product (NP, green), synthetic drug (SD, pink), and generated compound subgroups (Cheese, brown; STONED-SELFIES, blue). Horizontal dashed lines demarcate canonical Lipinski's rule of five upper thresholds: $\log P \leq 5$, molecular weight ≤ 500 Da, and polar surface area (TPSA) $\leq 140 \text{ \AA}^2$. Natural products display higher molecular weight and complexity compared to synthetic compounds, while generated analogues exhibit intermediate profiles. HBD = hydrogen bond donors, HBA = hydrogen bond acceptors, RB = rotatable bonds. Sample sizes: $n = 396$ NPs, 454 SDs, 40 481 generated compounds.




Figure 2: Bivariate consensus docking distributions across four *Plasmodium falciparum* targets: PfClpP (PDB: 4GM2), PfCRT (PDB: 6UKJ), PfDHFR-TS (PDB: 7F3Y), and PfATP4 (PDB: 9N10). 484 ligand centroids representing 65 856-molecule library screened per target. Compounds stratified into performance tiers: **EXCELLENT** (best score ≤ -7.0 kcal mol $^{-1}$ and DiffDock confidence ≥ 0.0 , red), **GOOD** (-7.0 kcal mol $^{-1}$ to -5.0 kcal mol $^{-1}$ and DiffDock ≥ -1.5 , orange), and **OTHERS** (grey). Y-axis: DiffDock confidence scores. PfCRT: 85.8 % hit rate; PfClpP: minimal enrichment. Classification thresholds calibrated against ChEMBL antimalarial actives (median -6.9 kcal mol $^{-1}$, IQR -7.4 kcal mol $^{-1}$ to -6.6 kcal mol $^{-1}$; $n = 37$).