

Supporting Information:

Computational Discovery of Antimalarial Candidates from African Natural Product Chemical Space

Myke Vital Sao Temgoua,^{*,†} Jean-Pierre Tchapet Njafa,[†] Penabei Samafou,[‡]
Wilfred Fon Mbacham,[¶] and Serge Guy Nana Engo[†]

[†]*Department of Physics, Faculty of Science, University of Yaoundé I, P.O. Box 812,
Yaoundé, Cameroon*

[‡]*Department of Medical Imaging and Radiation Sciences, Faculty of Medicine and Health
Sciences, Université de Sherbrooke, 3001, 12th Avenue Nord, Sherbrooke, QC J1H 5N4,
Canada*

[¶]*Department of Biochemistry, Faculty of Science, University of Yaoundé I, P.O. Box 812,
Yaoundé, Cameroon*

E-mail: myke-vital.sao@facsciences-uy1.cm

Introduction

This supplementary document details methodological protocols, scoring metrics, and dataset statistics supporting the main text (Temgoua *et al.*, 2026). The main manuscript describes a computational framework for antimalarial drug discovery that integrates African natural products with synthetic antimalarials, identifies antimalarial candidates (including select

multi-target hits) through dual-filter consensus docking, and prioritizes leads using multi-parameter optimization scoring.

The supplementary content elaborates the Multi-Parameter Optimization (MPO) scoring framework introduced in the main manuscript (Section 3.6), presenting mathematical formulations for the five desirability components and the polypharmacology bonus. Tabular summaries detail chemical space analysis, scaffold diversity, and drug-likeness characteristics of the 65 856-molecule hybrid library. Additional sections present generative model training and clustering validation, co-crystallized ligand classification for binding mode validation, computational resource comparisons, enrichment curves, UMAP projections of the latent space, and ScafVAE diversity benchmarks. Final tables provide target-specific docking outcomes, selectivity indices, and cytochrome P450 inhibition profiles for the 810 screened seed molecules with valid selectivity predictions against three *Plasmodium falciparum* structures: dihydrofolate reductase (DHFR, PDB: 7F3Y), chloroquine resistance transporter (CRT, 6UKJ), and ATPase 4 (ATP4, 9N10).

Calculation results referenced in this supplementary material are available from the project repository (https://github.com/NanaEngo/Malaria_codesV2) together with the analysis scripts in the `scripts/` directory.

Multi-parameter optimization scoring function

The MPO framework integrates five normalized desirability components into a composite score for hit prioritization. This approach extends the single-objective docking optimization strategy by simultaneously considering binding affinity, geometric confidence, drug-likeness, ADMET safety, and structural rule compliance. The methodology builds upon established multi-parameter optimization principles adapted for computational drug discovery applications. [S1,S2](#)

The total MPO score for compound i against target j is defined as:

$$\text{MPO}_{\text{total},i} = \sum_{k=1}^5 w_k \cdot S_{k,i} + \text{Polypharmacology Bonus}_i \quad (1)$$

where w_k are the component weights (Vina 35 %, DiffDock 25 %, QED 20 %, ADMET 15 %, Ro5 5 %) and $S_{k,i}$ are the normalized desirability scores scaled to $[0, 1]$. The polypharmacology bonus rewards compounds meeting the computational score threshold against multiple targets simultaneously; it does not establish biological polypharmacology.

Weight selection rationale

The component weights reflect established priorities in early-stage computational drug discovery. [S1,S2](#) The 35 % weight assigned to the Vina affinity score prioritizes the physics-inspired docking signal as the primary filter; this score is not a measured thermodynamic quantity. DiffDock geometric confidence (25 %) provides a complementary computational screening signal, [S3](#) with the combined 60 % weight for docking components reflecting the central role of target engagement in hit identification.

Drug-likeness (QED, 20 %) receives substantial weight to ensure compounds satisfy oral bioavailability criteria, [S4](#) while ADMET safety (15 %) addresses early-stage toxicity concerns. [S5](#) The modest 5 % weight for Lipinski violations reflects that Ro5 compliance is already captured partially within QED, and modern drug discovery increasingly tolerates single vi-

ulations for complex targets.^{S1} This weighting scheme balances potency, developability, and safety considerations typical of lead identification campaigns.^{S2}

Component 1: Vina binding affinity (S_{vina} , 35 %)

The Vina component is linearly scaled between the observed minimum and maximum binding affinities across all docked compounds:

$$S_{\text{vina},i} = \frac{\text{Vina}_i - \text{Vina}_{\min}}{\text{Vina}_{\max} - \text{Vina}_{\min}} \quad (2)$$

where Vina_i is the best binding affinity (kcal mol^{-1}) for compound i , and $\text{Vina}_{\min}$, $\text{Vina}_{\max}$ are the dataset bounds. More negative values yield higher desirability. This linear transformation preserves the rank ordering of compounds while mapping raw energies to a $[0, 1]$ scale compatible with other desirability components. Linear scaling was selected over exponential or power-law transformations to maintain proportional differences in binding affinity across the scoring range, consistent with the approximately linear relationship between binding free energy and experimental potency within the typical docking score range of $-10 \text{ kcal mol}^{-1}$ to -5 kcal mol^{-1} .

Component 2: DiffDock confidence (S_{diff} , 25 %)

The DiffDock confidence score is transformed via a logistic sigmoid function to map structural log-odds to probabilistic desirability:

$$S_{\text{diff},i} = \frac{1}{1 + e^{-c_i}} \quad (3)$$

where c_i is the DiffDock confidence score for compound i .^{S3} This transformation ensures that high-confidence poses ($c > 0$) receive scores approaching 1.0, while low-confidence poses ($c < -1.5$) receive scores approaching 0. The sigmoidal mapping was chosen over

linear scaling because DiffDock confidence scores follow a non-uniform distribution with a sharp transition around $c = 0$, and the sigmoid function naturally captures this threshold behavior while providing smooth gradation for intermediate confidence values.

Component 3: Quantitative estimate of drug-likeness (S_{qed} , 20 %)

The QED score is used directly as a desirability component without transformation:

$$S_{\text{qed},i} = \text{QED}_i \quad (4)$$

where $\text{QED}_i \in [0, 1]$ is the quantitative estimate of drug-likeness for compound i , computed from eight molecular properties using the RDKit implementation.^{S6} QED values > 0.70 are considered indicative of drug-like compounds.^{S4}

Component 4: ADMET safety profile (S_{admet} , 15 %)

The ADMET component aggregates 11 critical pharmacokinetic and safety parameters into a composite desirability index:

$$S_{\text{admet},i} = \frac{1}{11} \sum_{m=1}^{11} d_m(\text{ADMET}_{i,m}) \quad (5)$$

where d_m is the individual desirability function for ADMET endpoint m , predicted using ADMET-AI.^{S5} The 11 endpoints are:

- Oral bioavailability (human intestinal absorption, HIA)
- Aqueous solubility ($\log S$)
- Caco-2 permeability
- Blood-brain barrier penetration
- hERG liability (cardiotoxicity risk)

- AMES mutagenicity
- Hepatotoxicity (DILI)
- Synthetic accessibility (SA score)
- P-glycoprotein substrate status
- Cytochrome P450 inhibition risk (composite of CYP2C9, 2C19, 3A4, 2D6)
- Total clearance profile

Each endpoint is scored as a binary or continuous desirability function (1 = favourable, 0 = unfavourable), with thresholds derived from FDA-approved drug distributions.

Component 5: Lipinski’s rule of five penalty (P_{Ro5} , 5 %)

A penalty term proportional to formal Lipinski’s rule of five violations:^{S4}

$$P_{\text{Ro5},i} = 1 - 0.2 \cdot n_{\text{violations},i} \quad (6)$$

where $n_{\text{violations},i}$ is the number of Ro5 violations (0 to 4). The Ro5 criteria are: molecular weight > 500 Da, $\log P > 5$, hydrogen bond donors > 5, and hydrogen bond acceptors > 10. Penalties are: 0 violations = 1.0, 1 violation = 0.8, 2 violations = 0.6, 3 violations = 0.4, 4 violations = 0.2.

Polypharmacology bonus

Compounds exhibiting binding ($\text{MPO}_{\text{target}} \geq 0.50$) against multiple targets receive an additional bonus to reward multi-target activity:

$$\text{Polypharmacology Bonus}_i = 0.05 \cdot (n_{\text{targets},i} - 1) \quad \text{if } n_{\text{targets},i} \geq 2 \quad (7)$$

where $n_{\text{targets},i}$ is the number of targets against which compound i achieves $\text{MPO}_{\text{target}} \geq 0.50$. This bonus is capped at 0.10 (3 targets). The 0.05 increment per additional target represents a modest computational preference for multi-target profiles that may be relevant to resistance hypotheses; it does not demonstrate resistance circumvention or therapeutic advantage. The $\text{MPO} \geq 0.50$ threshold ensures that only compounds meeting minimum quality criteria across all scoring components qualify for the bonus, thereby distinguishing genuine multi-target candidates from non-selective binders with poor drug-likeness or safety profiles.

Centroid-based screening and MPO stratification

The MPO framework was applied within a centroid-based virtual screening workflow to enable efficient exploration of the 65 856-molecule library. Rather than exhaustively docking all compounds, the workflow proceeded as follows:

1. **K-Means clustering:** The library was partitioned into 484 clusters based on 64-dimensional VAE latent representations, with one centroid selected per cluster to represent its chemical neighborhood.
2. **Consensus docking:** The 484 centroids were docked against three *P. falciparum* targets (PfDHFR, PfCRT, and PfATP4) using dual-filter consensus scoring (AutoDock Vina + DiffDock). Centroids were classified as EXCELLENT, GOOD, or OTHERS based on binding affinity and geometric confidence thresholds. Failed records were excluded from consensus scoring and MPO calculation; target-specific denominators are reported with the corresponding results.
3. **MPO scoring:** 93 centroids passing consensus classification (EXCELLENT or GOOD) received full MPO evaluation across all five desirability components (Vina, DiffDock, QED, ADMET, Ro5).
4. **Optimization-worthy threshold:** 76 centroids achieved $\text{MPO} \geq 0.40$, defining the

optimization-worthy tier. Within this tier, 7 centroids achieved $\text{MPO} \geq 0.50$ (secondary hits), while 69 centroids fell in the 0.40 to 0.50 range (promising scaffolds).

5. **Scaffold expansion:** The 76 optimization-worthy centroids map to 53 unique clusters (multiple centroids can represent the same cluster because centroids are selected as representative molecules rather than cluster identifiers). Scaffold expansion from these 53 clusters yielded 40 481 molecules (mean cluster size: 763.8) for downstream lead optimization. Of these, 19 913 molecules satisfied the final synthetic accessibility and MPO criteria, defining the final set of synthesizable leads.

This centroid-based strategy reduced computational cost by 99.3 % compared to exhaustive library-scale docking (Table S33) while preserving chemical diversity through cluster-based sampling. The $\text{MPO} \geq 0.40$ threshold was selected to balance stringency with practical hit list size, ensuring that expanded clusters represent chemically tractable starting points for medicinal chemistry optimization.

Natural product relatedness filtering

To identify candidates retaining structural features of African natural products, NP-relatedness was assessed using Tanimoto similarity to 396 seed African NPs. Similarity was computed using Morgan fingerprints (radius 2, 2048 bits),^{S7} with compounds achieving maximum $\text{Tanimoto} \geq 0.4$ to any seed NP classified as NP-related. This threshold balances scaffold conservation with sufficient structural diversity for lead optimization.

Among the 76 optimization-worthy centroids, 13 (17.1 %) satisfied the NP-relatedness criterion ($\text{Tanimoto} \geq 0.4$), representing candidates that combine favorable multi-target profiles with natural product-inspired scaffolds. The remaining 63 centroids (82.9 %) represent structurally distinct analogues with $\text{Tanimoto} < 0.4$, confirming that generative expansion explored new chemical space while preserving drug-like characteristics.

This dual population supports two computational lead-optimization strategies: NP-

related candidates ($n = 13$) retain similarity to annotated seed scaffolds, whereas structurally distinct candidates ($n = 63$) expand the explored chemical space. Neither group establishes biological activity, toxicity, or resistance circumvention. The NP-relatedness filter was applied post-MPO scoring.

Framework limitations and sensitivity

Sensitivity methodology. The five MPO component scores comprise S_{vina} (Vina binding affinity), S_{diff} (DiffDock confidence), S_{qed} (QED drug-likeness), S_{admet} (ADMET composite), and P_{Ro5} (Lipinski penalty). Because individual Vina and DiffDock scores are not stored per molecule in the c6 library, S_{vina} and S_{diff} are back-calculated from the shared residual of the weighted MPO score after subtracting the QED, ADMET, and Ro5 contributions, then split proportionally (35:25 ratio). Weight perturbations of Vina and DiffDock therefore reallocate within this combined residual rather than varying the two components independently. QED, ADMET, and Ro5 perturbations reflect true independent variation.

The MPO framework represents a pragmatic balance between computational tractability and multi-objective optimization. Component weights were selected based on established drug discovery priorities^{S1,S4} rather than target-specific optimization, which would require experimental validation data currently unavailable for the majority of the library. All docking CSVs report provenance fields (target, exhaustiveness, grid version) so that every score can be traced to the receptor preparation and Vina search depth used. The framework’s primary limitation is the dataset-dependent nature of min-max scaling for Vina scores, which precludes direct score comparison across different compound libraries without renormalization.

Formal sensitivity analysis across 25 weight perturbations ($\pm 10\%$ to 20% per component) yielded mean Spearman rank correlation $\rho = 0.792$ for the top-1,000 compounds and mean Jaccard similarity of 0.548 for the top-20 hit list, indicating moderate sensitivity to weight choice (Table S19). When both stability criteria (top-1000 $\rho > 0.95$ and top-20 Jac-

card ≥ 0.70) are required, only 5 % of non-baseline perturbations passed (DiffDock 25 %; Vina, QED, ADMET, and Ro5 0 %), confirming that the ranking is moderately rather than strongly robust to weight choice; a substantial fraction of perturbations did not pass the stability test. The MPO ≥ 0.40 threshold for optimization-worthy centroids was selected to balance stringency with practical cluster expansion size, yielding 76 centroids that map to 53 unique clusters and expand to 40 481 molecules (19 913 final synthesizable leads) representing chemically tractable starting points for medicinal chemistry optimization.

Chemical space analysis and drug-likeness tables

The following tables summarize the chemical diversity, drug-likeness characteristics, and target-specific docking outcomes for the 65 856-molecule hybrid library.

Table S1: Top 10 most prevalent Bemis–Murcko scaffolds present in the 65 856-molecule generative hybrid library, highlighting the structural emphasis on privileged aromatic systems. *The library is dominated by benzene and quinoline-like privileged scaffolds; structural compatibility with antimalarial pharmacophores requires experimental or target-specific validation.*

Rank	Scaffold SMILES	Count	%	Description
1	c1ccccc1	5 537	8.4 %	Benzene ring
2	(acyclic)	3 610	5.5 %	Acyclic compounds
3	c1ccc2ncncc2c1	1 644	2.5 %	Quinazoline core
4	c1ccc(CCNCc2ccccc2)cc1	1 441	2.2 %	Diphenylethylamine
5	O=C(C=Cc1ccccc1)c1ccccc1	1 315	2.0 %	Chalcone scaffold
6	c1ccc2ncccc2c1	1 308	2.0 %	Quinoline core
7	O=c1cc[nH]c2ccccc12	1 176	1.8 %	Quinolin-2-one
8	c1ccc(Nc2ccnc3ccccc23)cc1	823	1.2 %	Phenylquinoline amine
9	O=c1ccc2ccccc2o1	798	1.2 %	Coumarin scaffold
10	O=C(Nc1ccccc1)c1ccccc1	787	1.2 %	Benzamide derivative
Summary: 20 702 unique scaffolds, 31.4 % diversity, Top 10 = 28.0 % of library				

Table S2: Distribution of predicted selectivity index proxies ($SI_{pred} = score_{eos7kpb}/(1 - DILI_{ADMET-AI})$) for 810 screened seed molecules with valid predictions. This computational proxy approximates antiparasitic selectivity over mammalian cytotoxicity; experimental $IC_{50}(HepG2)/IC_{50}(Pf\ NF54)$ determination is required for confirmation. *All screened molecules with valid SI predictions exceed the $SI_{pred} > 10$ proxy threshold; however, SI_{pred} is an unvalidated prioritization heuristic (sensitive to the DILI denominator, with an observed maximum of 106 773.8) and does not by itself establish a therapeutic window, which requires experimental $IC_{50}(HepG2)/IC_{50}(Pf\ NF54)$ determination.*

Metric	Value
Total molecules screened	810
Selectively antiparasitic ($SI > 10$)	100 % (810)
Mean selectivity index	993.8
Median selectivity index	283.5
Maximum selectivity index	106 773.8

Table S3: Cytochrome P450 (CYP) inhibition interaction risks for the 810 screened seed molecules with valid predictions, identifying moderate metabolic interaction probabilities (threshold > 0.5) relevant for formulation and drug-drug interactions. *Metabolic risk profiles are isoform-dependent, with CYP2C19 and CYP3A4 exhibiting the highest interaction probabilities.*

CYP Isoform	Mean	Median	High Inhibitors	%
CYP2C9	0.261	0.082	207	24.4 %
CYP2C19	0.424	0.346	356	41.9 %
CYP3A4	0.426	0.338	361	42.5 %
CYP2D6	0.170	0.001	138	16.2 %

Table S4: Stage-specific antimalarial activity predictions derived from the eos80ch model, indicating preferential predicted activity against asexual blood-stage parasites compared with transmission-blocking sexual gametocyte stages. *Identified candidates are predicted to primarily target the asexual blood stage, consistent with the clinical requirements for acute malaria treatment; these are computational predictions requiring experimental confirmation.*

Stage	Median	Mean	Interpretation
Asexual blood stage	0.567	0.647	Preferred predicted stage
Sexual stage (gametocytes)	0.244	0.502	Transmission-blocking potential

Table S5: Explained variance ratio across the first four principal components derived from PCA of nine standardized physicochemical descriptors (MW, logP, HBA, HBD, TPSA, QED, SA score, rotatable bonds, stereo centers). Components are ordered by decreasing explained variance ($PC1 \geq PC2 \geq PC3 \geq PC4$). The first three components capture 82.3 % of total structural variance. *Physicochemical diversity is highly compressible, with three components sufficing to capture the majority of library variance.*

Component	Variance	% Explained	Cumulative %
PC1	0.517	51.7 %	51.7 %
PC2	0.176	17.6 %	69.3 %
PC3	0.131	13.1 %	82.3 %
PC4	0.0686	6.86 %	89.19 %
PC1–3 total: 82.3 % of variance explained			

Table S6: Recovery rate of annotated African natural product (NP) scaffolds within the final expanded generative library, summarizing the retention of privileged bioactive cores throughout the STONED-SELFIES and Cheese API expansion steps. *The observed scaffold recovery is consistent with retention of some seed-framework features, but does not establish preservation of pharmacophore function.*

Metric	Value
Seed NP scaffolds	101
Library scaffolds	20 702
Recovered scaffolds	70
Recovery rate	69.3 %

Table S7: African natural product-related optimization candidates identified through centroid-based screening. NP-relatedness defined as Tanimoto ≥ 0.4 to 396 African NP seeds using Morgan fingerprints (radius 2, 2048 bits). *NP-related candidates span the full medicinal chemistry optimization spectrum, from secondary hits ready for lead optimization to promising scaffolds suitable for hit-to-lead development.*

Category	Total	NP-Related	Percentage	Mean Tanimoto
Secondary Hits (MPO ≥ 0.50)	7	2	28.6 %	0.52
Promising Scaffolds (MPO 0.40–0.50)	69	11	15.9 %	0.48
All Optimization Candidates	76	13	17.1 %	0.49

Table S8: Distribution metrics for the fraction of sp^3 hybridized carbons (f_{sp^3}), providing a quantitative measure of the three-dimensional structural complexity and saturation of the hybrid library compared to purely flat aromatic synthetic compounds. *The hybrid library maintains a balanced f_{sp^3} distribution, ensuring both structural complexity and synthetic tractability.*

Metric	Value
Mean f_{sp^3}	0.298
Median f_{sp^3}	0.286
Standard deviation	0.172

^aFinal synthesisable lead count (19,913) is taken from the canonical corrected-grid output ‘results/c6_primary_leads_synthesisable.csv’ and reflects SYBA > 0 plus the MPO $\geq$ 0.40 centroid filter; it supersedes the earlier 28,679 estimate.

Metric	Pass	Threshold	Rate
Full Library (65 856 molecules):			
Lipinski Ro5	61 975 / 65 856	All 4 criteria	94.1 %
Veber (rotatable bonds)	58 304 / 65 856	≤ 10	88.5 %
SYBA > 0	37 562 / 65 856	Synthetic accessibility	57.0 %
NPL > 0	28 014 / 65 856	NP-likeness	42.5 %
SA < 3.5	32 804 / 65 856	Synthetic tractability	49.8 %
QED > 0.70	24 144 / 65 856	Drug-likeness	36.7 %
Centroid-Based Screening Cascade:			
K-Means centroids docked	484 / 484	All clusters	100 %
Consensus classification pass	93 / 484	EXCELLENT or GOOD	19.2 %
Optimization-worthy (MPO $\geq$ 0.40)	76 / 93	Centroid MPO threshold	81.7 %
Secondary hits (MPO $\geq$ 0.50)	7 / 76	High MPO centroids	9.2 %
Scaffold Expansion from 53 Unique Clusters:			
Total molecules in expansion	40 481	Mean cluster size: 763.8	—
Final synthesizable leads ^a	19 913 / 40 481	SYBA > 0 + MPO filter	49.2 %

Table S10: Selected docking-predicted hits, including predicted multi-target candidates and single-target leads. These labels are computational predictions and do not establish biological activity.

Ligand	Target	Vina (kcal mol ⁻¹)	DiffDock Conf.	Notes
201	PfDHFR	−8.49	−0.30	Primary consensus lead
214–PfCRT	PfCRT	−8.33	−0.42	Named parent lead; MD system
87	PfDHFR	−8.38	−0.40	Primary consensus lead
438	PfATP4	−5.86	−0.95	Single GOOD hit
440	Multi-target	< −6.0	> −1.0	Polypharmacological
169	Multi-target	< −6.0	> −1.0	Polypharmacological
161	Multi-target	< −6.0	> −1.0	Polypharmacological
Predicted multi-target candidates in the expanded library: > 70 compounds meeting the declar				

Table S11: Top-20 compounds ranked by multi-parameter optimization (MPO) score. All compounds achieve perfect Lipinski compliance (4/4 criteria), high synthetic accessibility ($SA < 2.7$), and were generated via the STONED-SELFIES arm of the generative expansion. SMILES strings are provided for reproducibility. *MPO prioritization identifies candidates with favorable predicted drug-likeness, docking-score, and synthetic-accessibility profiles; these computational estimates require experimental validation.*

Rank	SMILES	MPO	QED	SA	SYBA	Source
1	<chem>Cc1ccc(CN2CCN(Cc3ccccc3)CC2)cc1O</chem>	0.829	0.939	1.72	109.01	STONED
2	<chem>Cc1ccccc1CN1CCN(Cc2cccc(O)c2)CC1</chem>	0.828	0.939	1.76	124.91	STONED
3	<chem>Cc1cc(CN2CCN(Cc3ccccc3)CC2)ccc1O</chem>	0.826	0.939	1.71	100.63	STONED
4	<chem>COc1ccc(CN2CCN(Cc3ccccc3C)CC2)cc1O</chem>	0.826	0.916	1.84	133.75	STONED
5	<chem>Oc1cccc(CN2CCN(Cc3ccc(Cl)cc3)CC2)c1</chem>	0.826	0.937	1.71	116.03	STONED
6	<chem>Cc1ccc(CN2CCN(Cc3ccc(O)cc3)CC2)c(C)c1</chem>	0.826	0.938	1.80	122.03	STONED
7	<chem>COc1ccc(CN2CCN(Cc3ccc(C)c(O)c3)CC2)cc1</chem>	0.826	0.916	1.80	119.14	STONED
8	<chem>Cc1ccc(CN2CCN(Cc3ccccc(O)c3)CC2)c(C)c1</chem>	0.825	0.938	1.85	129.53	STONED
9	<chem>Oc1ccc(CN2CCN(Cc3ccccc3)CC2)cc1Cl</chem>	0.825	0.937	1.72	113.80	STONED
10	<chem>Oc1ccc(CN2CCN(Cc3ccccc3Cl)CC2)cc1</chem>	0.824	0.937	1.71	122.53	STONED
11	<chem>COc1cccc(CN2CCN(Cc3ccc(C)c(O)c3)CC2)c1</chem>	0.824	0.916	1.85	120.90	STONED
12	<chem>Cc1ccc(CN2CCN(Cc3cc(O)ccc3C)CC2)cc1</chem>	0.823	0.938	1.82	106.23	STONED
13	<chem>Oc1cccc(CN2CCN(Cc3ccccc3Cl)CC2)c1</chem>	0.823	0.937	1.77	128.98	STONED
14	<chem>COc1ccc(CN2CCN(Cc3ccc(O)c(C)c3)CC2)cc1</chem>	0.823	0.916	1.79	110.76	STONED
15	<chem>Oc1ccc(CN2CCN(Cc3ccccc3)CC2)c(Cl)c1</chem>	0.823	0.937	1.76	113.45	STONED
16	<chem>COc1ccc(CN2CCN(Cc3ccccc(O)c3C)CC2)cc1</chem>	0.822	0.916	1.87	116.50	STONED
17	<chem>Oc1ccc(CN2CCN(Cc3ccccc(Cl)c3)CC2)cc1</chem>	0.822	0.937	1.71	115.24	STONED
18	<chem>COc1ccc([C@@H]2CC(=O)c3cc(O)cc(OC)cc3O2)cc1</chem>	0.822	0.943	2.62	41.36	STONED
19	<chem>Oc1ccc(CN2CCN(Cc3ccc(Cl)cc3)CC2)cc1</chem>	0.821	0.937	1.65	87.89	STONED
20	<chem>COc1cccc(CN2CCN(Cc3ccc(O)c(C)c3)CC2)c1</chem>	0.821	0.916	1.84	112.52	STONED
Summary: Mean MPO = 0.825 ± 0.003 ; Mean QED = 0.931 ± 0.010 ; Mean SA = 1.81 ± 0.20						
Drug-likeness: 20/20 compounds QED > 0.7; 20/20 compounds SA < 3.5; 20/20 compounds SYBA > 0						

Table S12: Constituent origins and corresponding molecule counts for the parallel generative expansion arms that assemble the final deduplicated 65 856-molecule hybrid antimalarial library. *The hybrid library combines similarity-based and perturbation-based expansion domains; the resulting computational profiles require experimental assessment.*

Source	File	Molecules	Description
REINVENT scaffolds	eos57bx_generated_mol2molscaf.csv	111 614	Novel scaffolds
COCONUT intermediates	eos6ost_np_mol_scaf_smi_filt.csv	172 750	NP scaffold expansion
Final hybrid library	eos80ch_malaria_final_activity.csv	65 856	Deduplicated, filtered library

Note: The counts for REINVENT scaffolds and COCONUT intermediates describe the raw generated candidate pools and intermediate datasets at intermediate stages of the workflow. They do not represent disjoint partitions of the library, and their sum does not equal the final deduplicated, filtered hybrid library size (65 856 molecules).

Table S13: Summary of cumulative explained variance for the Principal Component Analysis (PCA) used to characterize the 64-dimensional chemical space representation. *Low-dimensional representations capture the majority of structural variance, validating the use of compressed latent spaces for chemical space mapping.*

Metric	Value
PC1–PC3 cumulative variance	82.3 %
Number of molecular descriptors	200
Number of compounds analyzed	65 856

Table S14: Tanimoto structural similarity evaluation comparing generated analogues against seed African natural products. Novelty was assessed using Morgan fingerprints (radius 2, 2048 bits), defining structural novelty rigorously as a maximum Tanimoto coefficient < 0.4 relative to any foundational seed NP. *Over 92 % of the library represents structurally new chemical space, showing substantial separation from the evaluated natural-product reference set.*

Metric	Whole-Molecule	Scaffold-Only
Molecules analyzed	5 000 (random sample)	5 000 (same sample)
Seed NPs compared against	396	396
Novel molecules (Tanimoto < 0.4)	92.6 % (4 630)	—
Mean max Tanimoto	0.206	0.379
Std dev max Tanimoto	0.125	0.308
Median max Tanimoto	0.171	0.262
Whole-Molecule Tanimoto Distribution:		
< 0.2 (highly novel): 62.5 % (3 127)		
0.2 to 0.3 (novel): 25.2 % (1 258)		
0.3 to 0.4 (moderately novel): 4.9 % (245)		
0.4 to 0.5 (similar): 3.4 % (171)		
≥ 0.5 (very similar): 4.0 % (199)		

Table S15: Pearson correlation coefficient (r) analysis comparing DiffDock geometric confidence scores against AutoDock Vina affinity scores, characterizing the association between the two screening dimensions across the three evaluated targets. *The observed low-to-moderate associations support using Vina and DiffDock as complementary screening dimensions, while not proving independence or eliminating scoring bias.*

Target (PDB)	n ligands	Pearson r	p -value
PfDHFR (7F3Y)	280	0.327	2.11×10^{-8}
PfCRT (6UKJ)	64	0.129	3.08×10^{-1}
PfATP4 (9N10)	447	0.113	1.67×10^{-2}

Benchmark studies

Table S16: Enrichment analysis: MMV Malaria Box positive-control (Vina+DiffDock consensus) and DEKOIS PfDHFR independent docking (Vina only). *For two targets, the MMV comparison labels the top and bottom score strata and therefore measures retrodictive consistency rather than external discrimination. DEKOIS uses property-matched decoys docked with AutoDock Vina alone (no ML rescoring) as the independent baseline.*

Target (PDB)	ROC-AUC	EF1 %	EF5 %	EF10 %	BEDROC	PR-AUC	Actives	Decoys
PfDHFR-TS (7F3Y)	0.924	2.700	2.570	2.720	1.309	0.842	399	399
PfATP4 (9N10)	1.000	2.048	2.000	2.000	1.810	1.000	99	99
PfCRT (6UKJ)	0.971	1.110	1.110	1.110	9.434	0.992	359	40
<i>DEKOIS independent docking (Vina only, no ML rescoring):</i>								
PfDHFR-TS (7F3Y)	0.450	0.000	0.500	1.000	0.021	0.032	40	1 200

Notes: Top panel: MMV Malaria Box positive-control enrichment using Vina+DiffDock consensus scores; for PfCRT and PfATP4, the active/decoy labels are score-derived strata, so these values are retrodictive consistency checks rather than external discrimination. Bottom row: DEKOIS 2.0 PfDHFR independent docking using AutoDock Vina only (single-method, no ML rescoring). The near-random DEKOIS result (ROC-AUC 0.450) provides an external baseline for the limitations of docking-only discrimination; it motivates, but does not independently establish, the contribution of ML-based DiffDock rescoring because the consensus was not evaluated on the same DEKOIS panel. BEDROC (Boltzmann-Enhanced Discrimination of ROC, $\alpha = 20$).

Table S17: ChEMBL enrichment analysis: hit rates for ChEMBL-confirmed actives vs. inactives across three *P. falciparum* targets. *PfDHFR actives are enriched 5.43 $\times$ over inactives at the EXCELLENT threshold, supporting the use of stringent binding thresholds in this calibration set; PfATP4 shows no compounds reaching EXCELLENT; PfCRT’s high inactive hit rate suggests a permissive binding site.*

Target	PDB	Act.	Inact.	GOOD Hit Rate (%)		EXCELLENT Hit Rate (%)		Pass ($\geq 2\times$)	
				Act.	Inact.	Act.	Inact.	GOOD	EXC.
PfDHFR-TS	7F3Y	53	18	100.0 %	100.0 %	60.4 %	11.1 %	No	Yes
PfATP4	9N10	73	87	53.4 %	63.2 %	0.0 %	0.0 %	No	Yes
PfCRT	6UKJ	12	19	100.0 %	100.0 %	100.0 %	68.4 %	No	No

Notes: Actives ($IC_{50} < 1 \mu M$) vs. Inactives ($IC_{50} > 10 \mu M$) from ChEMBL. Hit rate thresholds reflect Vina scores (GOOD: $\leq -5.0 \text{ kcal mol}^{-1}$, EXCELLENT: $\leq -7.0 \text{ kcal mol}^{-1}$). Pass indicates ≥ 2 -fold enrichment of actives over inactives in the respective tier. The target set comprised PfDHFR, PfCRT, and PfATP4.

Integrated docking–RRS–polypharmacology panel

The following table presents a cross-source computational overlay of the target-anchored Vina panel with the P2 Set-C docking-RRS, PNS, and ACSI summaries. Candidate identifiers PP-01–PP-17 are linked by the source-bound molecular manifest using canonical SMILES verified against both the target-panel input manifest and the P2 Set-C source; the mapping and

Table S18: ADMET cross-referencing: Spearman correlation (ρ) between ADMET-AI and RDKit-based proxy descriptors (top-20 compounds). SwissADME estimates were not directly available (no public REST API); RDKit descriptors (ESOL logS and lipophilicity-based CYP heuristics) serve as SwissADME-equivalent proxies. ADMETlab 3.0 API returned 404 and was not available. *RDKit-based ADMET proxies provide a reproducible local alternative to web-based tools, with ESOL logS showing moderate correlation with ADMET-AI predictions.*

Endpoint	$\rho_{\text{AI vs Swiss}}$	$\rho_{\text{AI vs RDKit}}$	n	Disagreements	Notes
LogS (aqueous solubility)	-0.19	-0.19	20	0	ADMETlab 3.0 API unavailable
CYP3A4 inhibition	-0.29	-0.29	20	1	1 compound > 1 risk category apart

Note: ADMETLab 3.0 web API returned 404; values labelled “SwissADME” are RDKit-based proxies (ESOL logS and lipophilicity-based CYP heuristics). The identical ρ values for AI-vs-Swiss and AI-vs-RDKit columns confirm both use the same RDKit descriptors. Negative correlations reflect methodological differences in descriptor calculation. Correlations for hERG, BBB, and cLogP could not be computed due to insufficient proxy data (hERG/BBB required unavailable SwissADME descriptors; cLogP was not included in the ADMET-AI output used for this comparison).

source hashes are recorded explicitly in `results/v5_rrs_polypharm_integrated.provenance.json`.

The RRS column is docking-derived and target-specific, not MD-RRS; PNS and ACSI are prioritization descriptors, not biological measurements.

Interpretation boundary. The table is a cross-source computational overlay rather than an experimental or biological validation. The PP-to-source mapping is molecular-key verified against the integrated molecular manifest; the overlay is not a common-protocol recalculation. The legacy DiffDock consensus and the superseded four-target consensus are excluded.

The overlay CSV, mapping file, and SHA-256-bound source manifest are `results/v5_rrs_polypharm_integrated.provenance.json`, `results/v5_p2_candidate_mapping.csv`, and `results/v5_rrs_polypharm_integrated.provenance.json`.

Data availability

All raw calculation results referenced in this supplementary material are available in the project repository (https://github.com/NanaEngo/Malaria_codesV2):

- `results/c1_library_size_verification.txt`: Library size verification (65 856 molecules)

Table S19: MPO weight sensitivity analysis: 25 weight perturbations ($\pm 10\%$ to 20% per component). *The MPO ranking is moderately sensitive to weight choice: top-1000 Spearman ρ pass rate (> 0.95) is 28% and top-20 Jaccard pass rate (≥ 0.70) is 32% ; ADMET and QED weight perturbations produce the largest rank shifts.*

Weight Varied	Δ (%)	Vina (35 %)	DiffDock (25 %)	QED (20 %)	ADMET (15 %)	Ro5 (5 %)	ρ (top-1000)	Jaccard (top-20)	Pass Both
Vina weight perturbations									
vina	-20	0.188	0.313	0.250	0.188	0.063	0.821	0.290	No
vina	-10	0.278	0.278	0.222	0.167	0.056	0.938	0.482	No
vina	0	0.350	0.250	0.200	0.150	0.050	1.000	1.000	Yes
vina	10	0.409	0.227	0.182	0.136	0.045	0.943	0.818	No
vina	20	0.458	0.208	0.167	0.125	0.042	0.839	0.667	No
DiffDock weight perturbations									
diffdock	-20	0.438	0.063	0.250	0.188	0.063	0.890	0.379	No
diffdock	-10	0.389	0.167	0.222	0.167	0.056	0.966	0.538	Yes
diffdock	0	0.350	0.250	0.200	0.150	0.050	1.000	1.000	Yes
diffdock	10	0.318	0.318	0.182	0.136	0.045	0.968	0.905	Yes
diffdock	20	0.292	0.375	0.167	0.125	0.042	0.900	0.739	No
QED weight perturbations									
qed	-20	0.432	0.309	0.012	0.185	0.062	0.531	0.212	No
qed	-10	0.389	0.278	0.111	0.167	0.056	0.809	0.667	No
qed	0	0.350	0.250	0.200	0.150	0.050	1.000	1.000	Yes
qed	10	0.318	0.227	0.273	0.136	0.045	0.896	0.538	No
qed	20	0.292	0.208	0.333	0.125	0.042	0.798	0.429	No
ADMET weight perturbations									
admet	-20	0.407	0.291	0.233	0.012	0.058	0.258	0.000	No
admet	-10	0.389	0.278	0.222	0.056	0.056	0.418	0.053	No
admet	0	0.350	0.250	0.200	0.150	0.050	1.000	1.000	Yes
admet	10	0.318	0.227	0.182	0.227	0.045	0.622	0.333	No
admet	20	0.292	0.208	0.167	0.292	0.042	0.448	0.250	No
Ro5 weight perturbations									
ro5	-20	0.365	0.260	0.208	0.156	0.010	0.840	0.429	No
ro5	-10	0.365	0.260	0.208	0.156	0.010	0.840	0.429	No
ro5	0	0.350	0.250	0.200	0.150	0.050	1.000	1.000	Yes
ro5	10	0.318	0.227	0.182	0.136	0.136	0.643	0.333	No
ro5	20	0.292	0.208	0.167	0.125	0.208	0.435	0.212	No
Summary statistics							0.792	0.548	7/25

Note: Sensitivity analysis across 25 weight combinations ($\pm 10\%$ to 20% per component, others renormalized). Spearman ρ computed for top-1 000 compound ranking; Jaccard similarity for top-20 hit list. Pass thresholds: $\rho > 0.95$ AND Jaccard ≥ 0.70 . Only 7/25 combinations (28 %) pass both criteria. ADMET weight shows highest sensitivity (ρ drops to 0.26 at -20%). Vina and DiffDock weights are most stable ($\pm 10\%$ perturbations maintain $\rho > 0.93$).

Table S20: Sensitivity of generative expansion library size, scaffold diversity, and mean QED to CHEESE Tanimoto similarity threshold and STONED-SELFIES variant count per seed molecule. Baseline parameters (Tanimoto 0.7, 100 variants) are shown in **bold**. Values marked with † are approximate projections from the baseline and are not treated as measured results. *Generative expansion parameters are optimized to balance library size with structural diversity and drug-likeness.*

CHEESE Threshold	STONED-SELFIES Variants per Seed		
	50	100	200
0.6	68,849 / 13,634 / 0.587 [†]	80,999 / 15,149 / 0.582 [†]	93,149 / 15,906 / 0.572 [†]
0.7	55,975 / 12,394 / 0.602 [†]	65,856 / 13,772 / 0.597	75,730 / 14,460 / 0.587 [†]
0.8	38,622 / 10,535 / 0.607 [†]	45,438 / 11,706 / 0.602 [†]	52,254 / 12,291 / 0.592 [†]

Note: Analysis based on the actual generated library. Format: Library size / Unique scaffolds / Mean QED. The chosen parameters (0.7 threshold, 100 variants) produced 65,856 compounds with 13,772 unique scaffolds. CHEESE generated 29 805 compounds (394 scaffolds, QED 0.697). STONED-SELFIES generated 35 198 compounds (13 378 scaffolds, QED 0.497). † Estimated using first-order scaling from the baseline: $N_{\text{est}} = N_{\text{base}} \times f_{\text{STONED}} \times f_{\text{CHEESE}}$, where $f_{\text{STONED}} = V_{\text{variants}}/100$ and $f_{\text{CHEESE}} = (1.5 - T_{\text{CHEESE}})/0.8$ based on linear expansion limits.

Table S21: Cross-source computational overlay for the 17-candidate polypharmacology cohort. Vina values are target-anchored docking scores in kcal mol⁻¹; RRS is the corrected per-target docking-derived score; PNS is the polypharmacology network score; ACSI is the African Chemical Space Index. No column represents experimental affinity or efficacy.

Candidate	PfDHFR	PfCRT	PfClpP	PfATP4	Mean Vina	RRS class	RRS mean	PNS / ACSI
PP-01	-5.947	-6.483	-5.467	-6.361	-6.065	A*	87.78	4.45 / 0.459
PP-02	-5.562	-6.124	-5.907	-6.591	-6.046	A*	91.26	1.23 / 0.823
PP-03	-6.089	-6.633	-6.358	-7.311	-6.598	C	72.11	6.00 / 0.599
PP-04	-5.493	-5.329	-5.386	-6.086	-5.573	C	81.12	4.53 / 0.567
PP-05	-6.285	-6.412	-6.466	-6.378	-6.385	B	78.24	1.14 / 0.387
PP-06	-6.267	-7.907	-7.031	-7.285	-7.123	A*	92.37	1.04 / 0.589
PP-07	-6.226	-6.633	-6.226	-5.775	-6.215	B	75.01	4.48 / 0.603
PP-08	-4.935	-5.213	-5.052	-5.494	-5.174	C	70.80	4.36 / 0.605
PP-09	-5.693	-6.871	-5.613	-6.078	-6.064	B	74.18	4.20 / 0.601
PP-10	-6.346	-6.395	-6.016	-6.670	-6.357	B	80.49	4.48 / 0.551
PP-11	-6.179	-6.376	-6.583	-6.724	-6.466	A*	82.60	1.10 / 0.458
PP-12	-5.248	-6.357	-5.386	-4.635	-5.406	C	72.39	4.34 / 0.594
PP-13	-5.921	-5.685	-6.454	-7.565	-6.406	A*	104.82	2.32 / 0.597
PP-14	-4.863	-5.328	-5.047	-5.460	-5.175	D	69.53	4.36 / 0.743
PP-15	-6.045	-6.084	-6.366	-7.016	-6.378	A*	111.65	3.50 / 0.576
PP-16	-5.469	-6.202	-6.203	-6.451	-6.081	B	76.32	4.45 / 0.173
PP-17	-5.266	-5.116	-5.342	-5.360	-5.271	C	68.18	4.71 / 0.308

- `results/c2_per_source_breakdown.json`: Per-source breakdown of generative library
- `results/c3_selectivity_index.csv`: Selectivity index calculations (810/810 with $SI > 10$)
- `results/c4_cyp450_inhibition_profile.csv`: Cytochrome P450 inhibition predictions
- `results/c5_aqueous_solubility.csv`: Aqueous solubility analysis ($\log S$ values)
- `results/c6_primary_leads_synthesizable.csv`: Synthesizable primary leads (19 913 compounds)
- `results/c7_stage_specific_activity_summary.txt`: Stage-specific activity predictions
- `results/c8_pca_explained_variance.txt`: Principal component analysis results
- `results/c9_bemis_murcko_scaffolds.csv`: Bemis–Murcko scaffold analysis (20 702 unique scaffolds)
- `results/c10_scaffold_recovery_rate.txt`: Scaffold recovery rate (69.3 %)
- `results/c11_fsp3_per_source.csv`: Fraction of sp^3 carbons by source
- `results/c12_tanimoto_novelty_v2.csv`: Tanimoto novelty analysis
- `results/p1_threshold_calibration.csv`: Docking threshold calibration dataset
- `results/summary_consensus_*.csv`: Consensus docking results for target verification
- `results/p1_enrichment_chembl_benchmark.csv`: ChEMBL enrichment analysis results for the three evaluated targets

- `results/v5_rrs_polypharm_integrated.csv`: Source-hash-bound integration of the 17×4 Vina panel with P2 docking-RRS, PNS, and ACSI
- `results/v5_rrs_polypharm_integrated.provenance.json`: Overlay logic, source hashes, mapping status, and computational-status caveats
- `results/v5_p2_candidate_mapping.csv`: Source-bound PP identifier, rank, canonical-SMILES, and target-panel/P2 molecular mapping

The complete SMILES library (65 856 compounds) and available 3D structure archives (PDB/PDBQT formats) are versioned in the project repository. The source code and data used for the analyses are provided through the project repository (https://github.com/NanaEngo/Malaria_codesV2). A citable archival release will be provided separately.

Table S22: Target selectivity and safety profiling (top-20 MPO candidates). *Top candidates are predominantly single-target optimized molecules with favourable predicted selectivity; one candidate (rank 18) carries a predicted CYP3A4 inhibition flag.*

Classification	n	Criteria	Mean SI	Mean hERG	Mean CYP3A4	Flagged
Single-target (opt.)	20	$n_{\text{tgt}} = 1$, optimized for MPO	—	0.117	0.061	1

Note: Classification based on selectivity index (SI), hERG inhibition probability, and CYP3A4 inhibition risk from ADMET-AI. All 20 top candidates are single-target optimized molecules. Mean hERG = 0.117 (all < 0.5); mean CYP3A4 = 0.061, with 19 of 20 below 0.5 and rank 18 carrying a predicted CYP3A4 flag at 0.876. These are computational risk predictions requiring experimental follow-up. Of the > 70 multi-target candidates identified in the full expansion, 84% satisfied at least 3/5 ADMET risk-free criteria. All values are computational estimates requiring experimental validation.

Table S23: Detailed binding mode analysis for top-10 compounds ranked by MPO score. Key interactions represent computational predictions from AutoDock Vina docking poses and require experimental validation. Retrosynthetic feasibility assessed based on structural complexity and known reaction chemistry. *Structural analysis reveals high synthetic accessibility and conserved binding motifs across the top-ranked piperazine-based leads.*

Rank	SMILES	MPO	SA	QED
1	<chem>Cc1ccc(CN2CCN(Cc3ccccc3)CC2)cc1O</chem>	0.829	1.72	0.939
2	<chem>Cc1ccccc1CN1CCN(Cc2cccc(O)c2)CC1</chem>	0.828	1.76	0.939
3	<chem>Cc1cc(CN2CCN(Cc3ccccc3)CC2)ccc1O</chem>	0.826	1.71	0.939
4	<chem>COc1ccc(CN2CCN(Cc3ccccc3C)CC2)cc1O</chem>	0.826	1.84	0.916
5	<chem>Oc1cccc(CN2CCN(Cc3ccc(Cl)cc3)CC2)c1</chem>	0.826	1.71	0.937
6	<chem>Cc1ccc(CN2CCN(Cc3ccc(O)cc3)CC2)c(C)c1</chem>	0.826	1.80	0.938
7	<chem>COc1ccc(CN2CCN(Cc3ccc(C)c(O)c3)CC2)cc1</chem>	0.826	1.80	0.916
8	<chem>Cc1ccc(CN2CCN(Cc3cccc(O)c3)CC2)c(C)c1</chem>	0.825	1.85	0.938
9	<chem>Oc1ccc(CN2CCN(Cc3ccccc3)CC2)cc1Cl</chem>	0.825	1.72	0.937
10	<chem>Oc1ccc(CN2CCN(Cc3ccccc3Cl)CC2)cc1</chem>	0.824	1.71	0.937

Note: All compounds feature piperazine-based scaffolds with phenolic or methoxy substituents, suggesting potential for hydrogen bonding and hydrophobic interactions. SA scores < 2.0 indicate favorable predicted synthetic accessibility. QED scores > 0.9 indicate high predicted QED values. Retrosynthetic routes and 2D interaction diagrams are provided in Table S24 and Figure S8, respectively. Experimental validation of binding modes and synthetic routes is required before lead optimization.

Table S24: Predicted retrosynthetic routes for top-10 MPO-ranked candidates via ASKCOS (public API, MIT). Precursors and reaction types from template relevance model (Reaxys database).

Compound	Proposed Precursors	Reaction Type
T1 (route 1)	<chem>Cc1ccc(C=O)cc1O.c1ccc(CN2CCNCC2)cc1</chem>	Reductive amination (aldehyde + amine)
T1 (route 2)	<chem>C1CNCCN1.C=O.C=O.Cc1ccccc1O.c1ccccc1</chem>	Formylation / multicomponent
T1 (route 3)	<chem>COc1cc(CN2CCN(Cc3ccccc3)CC2)ccc1C</chem>	O-alkylation / ether formation
T2 (route 1)	<chem>Cc1ccccc1CN1CCNCC1.O=Cc1ccccc(O)c1</chem>	Reductive amination (aldehyde + amine)
T2 (route 2)	<chem>Cc1ccccc1C=O.Oc1ccccc(CN2CCNCC2)c1</chem>	Reductive amination (aldehyde + amine)
T2 (route 3)	<chem>Cc1ccccc1CN1CCNCC1.Oc1ccccc(CBr)c1</chem>	N-alkylation (alkyl bromide)
T3 (route 1)	<chem>Cc1cc(C=O)ccc1O.c1ccc(CN2CCNCC2)cc1</chem>	Reductive amination (aldehyde + amine)
T3 (route 2)	<chem>Cc1ccccc1O.c1ccc(CN2CCNCC2)cc1</chem>	C-N cross-coupling
T3 (route 3)	<chem>C1CNCCN1.Cc1cc(C=O)ccc1O.O=Cc1ccccc1</chem>	Reductive amination (two aldehydes + amine)
T4 (route 1)	<chem>COc1ccc(C=O)cc1O.Cc1ccccc1CN1CCNCC1</chem>	Reductive amination (aldehyde + amine)
T4 (route 2)	<chem>C1CNCCN1.C=O.C=O.COc1ccccc1O.Cc1ccccc1</chem>	Formylation / multicomponent
T4 (route 3)	<chem>COc1ccc(CO)cc1O.Cc1ccccc1CN1CCNCC1</chem>	N-alkylation (alcohol)
T5 (route 1)	<chem>Clc1ccc(CN2CCNCC2)cc1.O=Cc1ccccc(O)c1</chem>	Reductive amination (aldehyde + amine)
T5 (route 2)	<chem>O=Cc1ccc(Cl)cc1.Oc1ccccc(CN2CCNCC2)c1</chem>	Reductive amination (aldehyde + amine)
T5 (route 3)	<chem>Clc1ccc(CN2CCNCC2)cc1.Oc1ccccc(CBr)c1</chem>	N-alkylation (alkyl bromide)
T6 (route 1)	<chem>Cc1ccc(C=O)c(C)c1.Oc1ccc(CN2CCNCC2)cc1</chem>	Reductive amination (aldehyde + amine)
T6 (route 2)	<chem>C1CNCCN1.Cc1ccc(C=O)c(C)c1.O=Cc1ccc(O)cc1</chem>	Reductive amination (two aldehydes + amine)
T6 (route 3)	<chem>C1CNCCN1.Cc1ccc(CCl)c(C)c1.Oc1ccc(CCl)cc1</chem>	N-alkylation (alkyl chloride)
T7 (route 1)	<chem>COc1ccc(CN2CCNCC2)cc1.Cc1ccc(C=O)cc1O</chem>	Reductive amination (aldehyde + amine)
T7 (route 2)	<chem>COc1ccc(C=O)cc1.Cc1ccc(CN2CCNCC2)cc1O</chem>	Reductive amination (aldehyde + amine)
T7 (route 3)	<chem>COc1ccc(CN2CCNCC2)cc1.Cc1ccc(CCl)cc1O</chem>	N-alkylation (alkyl chloride)
T8 (route 1)	<chem>Cc1ccc(C=O)c(C)c1.Oc1ccccc(CN2CCNCC2)c1</chem>	Reductive amination (aldehyde + amine)
T8 (route 2)	<chem>Cc1ccc(CCl)c(C)c1.Oc1ccccc(CN2CCNCC2)c1</chem>	N-alkylation (alkyl chloride)
T8 (route 3)	<chem>Cc1ccc(CBr)c(C)c1.Oc1ccccc(CN2CCNCC2)c1</chem>	N-alkylation (alkyl bromide)
T9 (route 1)	<chem>O=Cc1ccc(O)c(Cl)c1.c1ccc(CN2CCNCC2)cc1</chem>	Reductive amination (aldehyde + amine)
T9 (route 2)	<chem>Oc1ccccc1Cl.c1ccc(CN2CCNCC2)cc1</chem>	C-N cross-coupling
T9 (route 3)	<chem>C1CNCCN1.C=O.C=O.Oc1ccccc1Cl.c1ccccc1</chem>	Formylation / multicomponent
T10 (route 1)	<chem>Clc1ccccc1CN1CCNCC1.O=Cc1ccc(O)cc1</chem>	Reductive amination (aldehyde + amine)
T10 (route 2)	<chem>O=Cc1ccccc1Cl.Oc1ccc(CN2CCNCC2)cc1</chem>	Reductive amination (aldehyde + amine)
T10 (route 3)	<chem>C1CNCCN1.O=Cc1ccc(O)cc1.O=Cc1ccccc1Cl</chem>	Reductive amination (two aldehydes + amine)

SM Figure S15: Top-10 Polypharmacological Candidates




Figure S1: 2D structure diagrams of the top-10 MPO-ranked candidates. All compounds feature a piperazine linker connecting two substituted aryl rings, with hydrogen bond donors (phenolic -OH) and acceptors (methoxy -OCH₃, piperazine N). Molecular weights range 340 Da to 430 Da, consistent with oral drug-like space. The top-10 MPO subset is predominantly single-target optimized (see Supplementary Table S22).

Table S25: Benjamini–Hochberg FDR-corrected pairwise comparisons of antimalarial activity scores across compound groups (NP, SD, Cheese, STONED). *Multiple comparison testing found no statistically significant difference in the analysed activity scores between generated analogues and the foundational natural-product and synthetic-drug groups; this does not establish equivalent experimental potency.*

Model	Groups	p (raw)	p (adj)	Sig.
eos7yti	NP vs SD	> 0.05	> 0.05	n.s.
eos7yti	NP vs Cheese	> 0.05	> 0.05	n.s.
eos7yti	NP vs STONED	> 0.05	> 0.05	n.s.
eos7yti	SD vs Cheese	> 0.05	> 0.05	n.s.
eos7yti	SD vs STONED	> 0.05	> 0.05	n.s.
eos7yti	Cheese vs STONED	> 0.05	> 0.05	n.s.
eos7kpb	NP vs SD	> 0.05	> 0.05	n.s.
eos7kpb	NP vs Cheese	> 0.05	> 0.05	n.s.
eos7kpb	NP vs STONED	> 0.05	> 0.05	n.s.
eos7kpb	SD vs Cheese	> 0.05	> 0.05	n.s.
eos7kpb	SD vs STONED	> 0.05	> 0.05	n.s.
eos7kpb	Cheese vs STONED	> 0.05	> 0.05	n.s.
eos80ch	NP vs SD	> 0.05	> 0.05	n.s.
eos80ch	NP vs Cheese	> 0.05	> 0.05	n.s.
eos80ch	NP vs STONED	> 0.05	> 0.05	n.s.
eos80ch	SD vs Cheese	> 0.05	> 0.05	n.s.
eos80ch	SD vs STONED	> 0.05	> 0.05	n.s.
eos80ch	Cheese vs STONED	> 0.05	> 0.05	n.s.
eos4zfy	NP vs SD	> 0.05	> 0.05	n.s.
eos4zfy	NP vs Cheese	> 0.05	> 0.05	n.s.
eos4zfy	NP vs STONED	> 0.05	> 0.05	n.s.
eos4zfy	SD vs Cheese	> 0.05	> 0.05	n.s.
eos4zfy	SD vs STONED	> 0.05	> 0.05	n.s.
eos4zfy	Cheese vs STONED	> 0.05	> 0.05	n.s.

Note: 24 pairwise Mann–Whitney U tests (4 models $\times$ 6 group pairs). Benjamini–Hochberg FDR correction applied to control false discovery rate at $\alpha = 0.05$. All 24 comparisons: adjusted $p > 0.05$. Given the large sample size of the expansion groups ($N \approx 32\,000$), which creates excessive statistical power, the rank-biserial correlation ($r < 0.10$, indicating negligible effect size) serves as the primary metric to confirm the absence of meaningful differences between the generated groups and the seeds. Group sample sizes: NP ($n = 396$), SD ($n = 454$), Cheese ($\approx 32\,000$), STONED ($\approx 33\,000$).

Co-crystallized ligand classification

To address potential reviewer concerns regarding the biological relevance of co-crystallized ligands in the target protein structures used for docking, we employed LigandExplorer v2.1 (DPTech, 2026) to automatically classify ligands extracted from the four *P. falciparum* crystal structures (PDB: 7F3Y, 6UKJ, and 9N10). LigandExplorer uses a graph neural network (GNN) backend with pocket-aware graph construction to classify molecules into biological types (DNA, RNA, peptide, glycan, lipid, organic, ion) and assess their relevance as biologically meaningful ligands rather than crystallization artifacts.^{S8}

Table S26: LigandExplorer v2.1 classification of co-crystallized ligands extracted from the four *P. falciparum* target structures used in this study. Classification performed using the GNN backend (PyTorch 2.12.1) with whole-structure search mode and default parameters. *All detected co-crystallized ligands are classified by the software as DNA, RNA, peptide, or organic covalent; this software classification does not independently establish biological relevance or exclude crystallization artifacts.*

PDB	Target	Ligand	GNN Classification	Chain	Relevance
7F3Y	PfDHFR-TS	UMP	DNA-like ligand	A, B	Biological
		MTX	Peptide-like ligand	A ($\times 2$), B	Biological
		NDP	RNA-like ligand	A, B	Biological
		GOL	Organic (crystallization)	A, B	Artifact
6UKJ	PfCRT	Y01	Organic covalent ligand	A	Biological
		LEU	Peptide covalent ligand	A	Biological
9N10	PfATP4	—	No ligands detected	—	—

Note: LigandExplorer v2.1 (<https://github.com/dptech-corp/ligandexplorer>) with GNN backend. UMP (uridine monophosphate) is a nucleotide cofactor of thymidylate synthase; MTX (methotrexate) is a classical DHFR inhibitor; NDP (NADPH) is a nucleotide cofactor; Y01 is a chloroquine analogue; LEU is a leucine residue. GOL (glycerol) is a common crystallization agent correctly identified as a non-biological artifact. 9N10 stores ligand records in PDB format variant that differs from the expected LigandExplorer input format; manual inspection confirms their co-crystallized ligands are present in the source structures.

Model training and clustering validation

Variational autoencoder (VAE) training

Two VAE latent space dimensions were evaluated (32D and 64D). Training converged successfully for both dimensionalities (Figure S2, Table S27).

Table S27: VAE latent space comparison: 32D vs 64D. The ScafVAE GuacaMol distribution-learning KL divergence score (0.996) provides an external reference benchmark. The per-dimension VAE training metrics (validation loss, reconstruction loss, KL divergence, reconstruction accuracy) and the latent-space clustering indices (Silhouette, Calinski–Harabasz, Davies–Bouldin) were not retained post-hoc and are reported as N/A; the 64-D latent space was selected for centroid identification on the basis of the clustering-algorithm benchmark (Supplementary Table S29) and downstream centroid selection, not on the metrics listed here.

Metric	32D	64D
Validation Loss	N/A	N/A
Reconstruction Loss	N/A	N/A
KL Divergence	N/A	N/A
Reconstruction Accuracy (%)	N/A	N/A
ScafVAE Baseline KL Div. Score ^a S9	0.996	
Silhouette Score	N/A	N/A
Calinski-Harabasz Index	N/A	N/A
Davies-Bouldin Index	N/A	N/A

^a GuacaMol benchmark distribution-learning KL divergence score.

Evaluation of clustering strategies and dimensionality

Comparison of 32-dimensional versus 64-dimensional latent spaces revealed trade-offs in clustering performance (Figure S3). The 64-dimensional space was selected for centroid selection based on comparative clustering analysis and the clustering-quality figure; the retained summary does not provide numeric indices sufficient to claim metric-optimality.

To partition the library into chemical families, KMeans, BIRCH, and Jarvis-Patrick clustering algorithms were optimized via Optuna^{S11} across 32D and 64D VAE latent spaces. The 64-dimensional KMeans configuration ($n = 484$) was selected as a practical balance of




Figure S2: VAE training history for 32D (top) and 64D (bottom) latent spaces. Both models converged successfully with stable validation loss in the final epochs. Training curves show reconstruction loss, KL divergence, and validation accuracy over 100 epochs. *VAE training successfully converges for both 32- and 64-dimensional spaces, with the latter selected for comparative chemical-space resolution; the retained summary does not establish superiority from a complete set of numeric validation indices.*

Table S28: Comparison of generative model novelty and diversity metrics with the ScafVAE state-of-the-art. ScafVAE scores are published GuacaMol distribution-learning benchmarks.^{S10} Our metrics are computed on the African NP / synthetic hybrid library against 396 seed African NPs (Morgan fingerprint Tanimoto at threshold ≥ 0.4). Direct numerical comparison is limited by different evaluation frameworks (GuacaMol distribution learning vs. target-aware virtual screening).

Metric	This work	ScafVAE ^{S9}
Benchmark framework	NP-anchored virtual screening	GuacaMol distribution learning
Training dataset	ANPDB + DrugBank (NP-focused)	ChEMBL / ZINC (drug-like)
Validity / Valid SMILES	99.93 %	0.987 to 0.997
Novelty (unreachable fraction)	92.6 % ($T_c < 0.4$ to seeds)	1.000
Scaffold diversity (unique Bemis-Murcko)	31.4 %	—
Scaffold recovery from seeds	69.3 %	—
GuacaMol KLD	—	0.959
GuacaMol FCD	—	0.338




Figure S3: Clustering quality metrics comparison between 32D and 64D latent spaces. Left: Silhouette score (higher is better, >0.5 = good, >0.25 = lower bound acceptable). Middle: Calinski-Harabasz index (higher is better). Right: Davies-Bouldin index (lower is better). The 64D space was selected for centroid selection based on the clustering metrics shown. *The 64-dimensional latent space was selected as the representation for chemical clustering and centroid selection based on these comparative metrics.*

granularity and cohesion (Table S29); the retained summary does not contain the numeric internal indices needed to claim metric-optimality. Jarvis–Patrick yielded fragmented structures (42 917–44 040 clusters), while BIRCH produced heterogeneous partitions with poor Silhouette scores.

Table S29: Comparative performance metrics for KMeans, BIRCH, and Jarvis-Patrick clustering algorithms mapped across 32-dimensional and 64-dimensional VAE latent spaces for the complete hybrid library ($N = 65\,856$). Hyperparameters were systematically optimized via Optuna. CH = Calinski–Harabasz index, DB = Davies–Bouldin index. KMeans on the 64-D latent space ($n = 484$) was selected for centroid identification after benchmarking against BIRCH (heterogeneous partitions, poor Silhouette) and Jarvis–Patrick (fragmented, 42 917 to 44 040 clusters); the selected configuration’s internal indices are not retained in this summary, and selection followed the comparative benchmarking and practical cluster granularity rather than a metric shown as optimal here.

Algorithm	Dim	Params (k, s, thr)	Sil.	CH	DB	Clusters (Min–Max)	Quality
Jarvis–Patrick	32	$k = 151, s = 138$	N/A	N/A	N/A	42 917 (1 to 5)	Acceptable
	64	$k = 181, s = 171$	0.274	566	0.376	44 040 (1 to 2)	Acceptable
BIRCH	32	$n = 490, thr = 0.11$	0.122	703	1.585	490 (13 to 820)	Poor
	64	$n = 489, thr = 0.10$	0.162	1 187	1.405	489 (11 to 858)	Poor
KMeans	32	$n = 499$	0.182	875	1.392	499 (12 to 454)	Poor
	64	$n = 484$	N/A	N/A	N/A	484 (11 to 519)	Acceptable*

Quality assessment: Silhouette score interpretation: > 0.50 = good separation, 0.25 to 0.50 = acceptable (lower bound for continuous chemical space), < 0.25 = poor. Calinski–Harabasz: higher = better inter-cluster separation. Davies–Bouldin: lower = better cluster compactness. The selected 64-D KMeans configuration ($n = 484$) was adopted for centroid selection based on the comparative benchmarking across algorithms and its practical cluster granularity (member range 11 to 519); its internal Silhouette, Calinski–Harabasz, and Davies–Bouldin values are not retained in this summary table.

Docking protocol validation (redocking)

Redocking assessment was performed on co-crystallized ligands from PfDHFR-TS (PDB: 7F3Y) and PfCRT (PDB: 6UKJ). RMSD values between predicted and crystal poses were calculated using RDKit’s molecular alignment algorithm. An RMSD threshold of $2.0\text{ \AA}$ was used to define reference-pose reproduction.^{S12} Among alignable ligands, 5/5 were below this

threshold (5/10 overall, with 5 ligands excluded due to SDF atom-ordering mismatches), with mean RMSD of $1.33 \text{ \AA} \pm 0.43 \text{ \AA}$ (Table S30, Figure S4). This limited check does not establish prospective docking accuracy. Binding affinities of non-aligned ligands (Vina $-7.01 \text{ kcal mol}^{-1}$ to $-7.69 \text{ kcal mol}^{-1}$) provide score context but do not independently establish correct active-site placement.

Table S30: Redocking assessment for co-crystallized ligands. *The protocol reproduced the reference poses for the evaluated alignable ligands at $\text{RMSD} < 2.0 \text{ \AA}$; this limited computational check does not establish prospective docking accuracy.*

Target	Ligand	Vina (kcal mol^{-1})	Vinardo (kcal mol^{-1})	RMSD ($\text{\AA}$)	Pass ($< 2.0 \text{ \AA}$)
7F3Y	UMP	-5.53	-4.14	1.005	Yes
7F3Y	GOL	-3.40	-3.14	1.002	Yes
7F3Y	UMP	-5.59	-4.16	1.882	Yes
7F3Y	GOL	-3.43	-2.97	1.029	Yes
6UKJ	Y01	-9.08	-6.40	1.713	Yes
7F3Y	NDP	-7.20	-4.53	N/A	No ^a
7F3Y	MTX	-7.01	-4.79	$\sim 30 \text{ \AA}$	No ^a
7F3Y	NDP	-7.69	-4.12	N/A	No ^a
7F3Y	MTX	-7.03	-4.68	$\sim 30 \text{ \AA}$	No ^a
7F3Y	MTX	-7.22	-4.73	$\sim 30 \text{ \AA}$	No ^a

Success Rate: 5/10 (50.0 %) overall; 5/5 (100.0 %) for alignable ligands

^aRMSD could not be computed via standard protocol due to SDF atom ordering mismatch; cross-analysis with DiffDock is consistent with ligand dissociation and an RMSD of approximately $30 \text{ \AA}$; this result illustrates the inability of rigid-receptor docking to reproduce the flexible DHFR active-site MTX binding mode.

Note: Binding affinities for non-aligned ligands provide score context only; they do not independently establish correct active-site placement.

Dual-filter consensus validation

Comparison of single-method versus dual-filter consensus screening characterized the effect of combining two computational scoring signals (Table S31, Figure S5).¹ For Vina-only classification, compounds with scores $\leq -7.0 \text{ kcal mol}^{-1}$ were classified as EXCEL-

¹Hit rates and reduction percentages in Table S31 were computed with EXCELLENT $\leq -7.0 \text{ kcal mol}^{-1}$, GOOD $-7.0 \text{ kcal mol}^{-1}$ to $-5.0 \text{ kcal mol}^{-1}$ Vina-only thresholds, consistent with the main text calibration.




Figure S4: Redocking validation results showing RMSD between predicted and crystal poses for co-crystallized ligands from PfDHFR-TS (7F3Y) and PfCRT (6UKJ). All successfully aligned ligands achieved RMSD < 2.0 Å (red dashed line), indicating good pose reproduction for the evaluated alignable ligands. *This computational redocking check does not establish prospective docking accuracy for new compounds.*

LENT, $-7.0 \text{ kcal mol}^{-1}$ to $-5.0 \text{ kcal mol}^{-1}$ as GOOD, and $> -5.0 \text{ kcal mol}^{-1}$ as OTHERS. For DiffDock-only, confidence scores ≥ 0.0 were classified as HIGH, -1.5 a to 0.0 as MODERATE, and < -1.5 as LOW. Consensus classification required passing both filters. For PfCRT, consensus filtering achieved a 33.0 % reduction (from 97.6 % to 64.5 %), while PfATP4 showed the largest reduction of 45.2 % (from 63.9 % to 18.6 %). Correlations between Vina scores and DiffDock confidence were low to moderate (Figure S6), consistent with treating the scores as complementary screening signals; low correlation alone does not establish independence or eliminate shared scoring bias.

Table S31: Consensus hit rates from Vina-only, DiffDock-only, and dual-filter consensus screening. *The dual filter narrows the computational hit set; an independent inactive panel is required to quantify any false-positive reduction.*

Target	Vina Only (% hits)	DiffDock Only (% hits)	Consensus (% hits)	Reduction (%)
PfDHFR-TS	94.9	87.1	82.0	12.9
PfCRT	97.6	66.3	64.5	33.0
PfATP4	63.9	30.8	18.6	45.2




Figure S5: Hit-rate comparison between single-method and dual-filter consensus screening. Consensus filtering reduces the fraction of compounds passing the declared thresholds across targets. This is a screening-set hit-rate change, not a measured false-positive reduction, because an independent inactive panel was not available for all targets. *Consensus filtering narrows the computational candidate pool, particularly for challenging targets such as PfATP₄.*



Figure S6: Correlation between Vina scores and DiffDock confidence for the three evaluated targets. Low to moderate correlations are consistent with the two methods providing complementary screening signals, but do not establish statistical independence or validate pose quality. Threshold lines indicate classification boundaries for EXCELLENT/GOOD (Vina) and HIGH/MODERATE (DiffDock) categories. Colors indicate consensus classification: green (EXCELLENT), orange (GOOD), gray (OTHERS). *An independent inactive panel and pose-level validation would be required to quantify complementarity and predictive value.*

MMV Malaria Box positive control validation

Table S32: MMV Malaria Box positive control validation. Hit rates (HIGH + MEDIUM tiers) for 399 compounds with confirmed antimalarial activity against *P. falciparum* 3D7 ($IC_{50} < 10 \mu M$).

Target	N	Hit Rate (H+M)	HIGH	MEDIUM	Best Vina (kcal mol ⁻¹)
PfDHFR	399	35.1 %	132	8	-10.12
PfCRT	399	90.7 %	359	3	-13.37
PfATP4	184	47.3 %	0	87	-6.94

All 399 MMV Malaria Box compounds have confirmed antimalarial activity. PfATP4 shows $N = 184$ due to the 15-rotatable-bond filter. $H+M = \text{HIGH (Vina} \leq -7.0 \text{ kcal mol}^{-1}) + \text{MEDIUM} (-7.0 \text{ kcal mol}^{-1} \text{ to } -5.0 \text{ kcal mol}^{-1})$ tiers. One compound (1/400) failed PDBQT conversion.

Computational resource comparison

Table S33: Computational resource comparison: exhaustive screening versus centroid-based sampling for the 65 856-molecule library across three *P. falciparum* targets. *Centroid-based sampling reduces the number of docking representatives by more than 99 %; this computational reduction does not establish prospective hit-identification accuracy.*

Metric	Exhaustive	Centroid-Based
Molecules Screened	65 856	484
Docking Runs (3 targets)	197 568	1 452
CPU-Hours (estimated) ^a	70 248 to 175 620	517 to 1 291
Cost Reduction	:	99.3 %
ROC-AUC (PfDHFR)	Unknown ^b	0.924 [0.895 to 0.951]
ROC-AUC (PfATP4)	Unknown ^b	1.000 [1.000 to 1.000]
ROC-AUC (PfCRT)	Unknown ^b	0.971 [0.949 to 0.989]
Accessibility	Limited ^c	High

^a Based on AutoDock Vina (exhaustiveness = 64, 32 CPU cores), 16 to 40 CPU-minutes per docking run.

^b Exhaustive screening not performed due to computational constraints.

^c Requires > 70 000 CPU-hours, prohibitive for resource-limited settings.

Table S34: Virtual screening results for cluster centroids across three targets. 484 centroids representing 65 856 compounds were screened per target.

Target	PDB	Total	Excellent (%)	Good (%)	Hit Rate (%)	Best Vina (kcal mol ⁻¹)
PfATP4	9N10	484	0.0	0.4	0.4	-7.11
PfCRT	6UKJ	484	20.6	65.2	85.8	-12.27
PfDHFR-TS	7F3Y	484	1.8	35.9	37.7	-9.67

Virtual screening results

Enrichment curves

Figure S7: Enrichment curves for the three evaluated targets. Each panel shows: (A) ROC curve with AUC value, (B) Precision-Recall curve, (C) Enrichment factor vs database fraction, and (D) Score distribution for actives (green) and decoys (red). All metrics calculated from actual docking results with bootstrap confidence intervals (1 000 resamples). PfDHFR assessed against DEKOIS benchmark; other targets assessed using MMV Malaria Box compounds.

Lead interaction diagrams




Figure S8: 2D protein-ligand interaction diagrams for named parent leads. (a) Ligand 201 bound to *PfDHFR*-TS (7F3Y; Vina -8.49 kcal mol⁻¹), showing predicted hydrogen-bond interactions with Asp54 and Arg57. (b) Ligand 214 bound to *PfCRT* (6UKJ), showing high DiffDock geometric confidence. (c) Ligand 438 bound to *PfATP4* (9N10; Vina -5.86 kcal mol⁻¹), the sole GOOD hit for this target. Diagrams generated with ChimeraX. These parent systems were analyzed separately from the polypharmacology cohort.

Supplementary figures




Figure S9: UMAP projection of the 64-dimensional VAE latent space, colored by compound source. Natural products (NP, green, $n = 396$), synthetic drugs (SD, red, $n = 454$), and computationally generated analogues (Generated, blue, $n = 65\,856$) occupy overlapping but distinguishable regions of chemical space. The continuous distribution is consistent with integration of natural-product scaffolds and synthetic drug-like properties during generative expansion. NP compounds cluster toward negative UMAP 1 values (centroid: $-0.177, 0.191$), while SD and generated compounds occupy more positive regions (SD centroid: $0.082, 0.363$; Generated centroid: $0.206, 0.239$), reflecting scaffold diversification in the Cheese similarity-search and STONED-SELFIES perturbation arms. UMAP parameters: `n_neighbors = 15`, `min_dist = 0.1`. *The projection is descriptive and does not by itself establish successful biological integration or activity.*




Figure S10: Physicochemical property distributions extracted from the ten most populated subfamilies defined by 64-dimensional KMeans clustering, demonstrating distinct, chemically coherent gradients within the overall library. White internal markers denote cluster medians. *Clustering identifies structurally coherent molecular families with distinct physicochemical gradients, enabling target-specific scaffold selection.*

Enrichment validation - supplementary information

Benchmark datasets and computational design

Data sources

The validation metrics were calculated from the docking datasets and protocols described below. The corresponding input files and processing scripts are identified to facilitate reproduction:

MMV Malaria Box:

- **Source:** Medicines for Malaria Venture (MMV)
- **Compounds:** 400 known antimalarials
- **Activity:** All confirmed active against *P. falciparum* 3D7 strain
- **IC₅₀:** All < 10 μ M
- **Reference:** Spangenberg et al. (2013), PLoS ONE
- **Access:** Publicly available from MMV
- **File:** data/raw/MalariaBox400compoundsDec2014.csv

DEKOIS 2.0 Benchmark:

- **Source:** DEKOIS 2.0 database (Bauer et al., 2013)
- **Target:** PfDHFR (DHFR ligands and decoys)
- **Actives:** 40 experimentally confirmed
- **Decoys:** 1,200 carefully matched (30:1 ratio)
- **Design:** Similar physicochemical properties, different topology

- **Reference:** Bauer et al. (2013), J. Chem. Inf. Model.
- **Files:** data/external/DHFR_ligands.smi, data/external/DHFR_decoys.smi

AutoDock Vina/Vinardo Results:

- **Source:** Our docking protocol
- **Targets:** 3 *P. falciparum* proteins (PfDHFR, PfCRT, PfATP4)
- **Method:** AutoDock Vina 1.2.7 with Vina and Vinardo scoring
- **Parameters:** Exhaustiveness = 64, CPU cores = 32, time per run = 16–40 CPU-minutes
- **Files:** Docking/Docking_[PDB]/mmv_results_consensus/summary_consensus_[PDB].csv

DiffDock Results:

- **Source:** Our docking protocol
- **Method:** DiffDock (Corso et al., 2023)
- **Poses:** 40 per ligand-target pair
- **Confidence:** Geometric confidence scores
- **Files:** Docking/results_inference_mmv/[PDB]_mol_[N]/rank1_confidence-[score].sdf

Benchmark design

The benchmark design combined a physicochemically matched decoy set, complementary enrichment metrics, and bootstrap confidence intervals:

1. **Matched Decoys:** Use carefully designed decoys (DEKOIS) to avoid trivial separation

2. **Multiple Metrics:** Calculate complementary metrics (ROC-AUC, EF, BEDROC)
3. **Statistical Confidence:** Bootstrap confidence intervals (1 000 resamples)
4. **Traceable analysis:** Record the data sources, processing steps, and analysis parameters

Analysis reproducibility

All Vina and DiffDock values were obtained from the docking runs described in the Methods. Classification thresholds were specified from the reference activity distributions and applied without post hoc score adjustment. The source datasets, analysis scripts, and derived result files are listed below to support independent reproduction.

Validation script: `scripts/revision/data_driven_enrichment_validation.py`

Results files:

- `results/data_driven_enrichment_validation_results.csv`
- `results/validation_dataset_PfDHFR_7F3Y.csv`
- `results/validation_dataset_PfATP4_9N10.csv`
- `results/validation_dataset_PfCRT_6UKJ.csv`

Bootstrap confidence interval methodology

Bootstrap confidence intervals were calculated using 1 000 resamples with replacement. For each resample:

1. Randomly sample N compounds (with replacement) from the validation set, where N is the original dataset size
2. Calculate ROC-AUC, EF5 %, EF10 %, and BEDROC for the resample
3. Store the metric values

The 95 % confidence interval was determined as the 2.5th and 97.5th percentiles of the bootstrap distribution. Narrow confidence intervals indicate consistent performance across different dataset compositions.

Enrichment metric definitions

ROC-AUC (Receiver Operating Characteristic - Area Under Curve): Measures the probability that a randomly chosen active compound ranks higher than a randomly chosen decoy. Range: $[0, 1]$, where 1.0 indicates perfect separation and 0.5 indicates random ranking. ROC-AUC is threshold-independent and provides an overall assessment of ranking quality.

Enrichment Factor at 5 % (EF5 %): Quantifies how many times more actives are found in the top 5 % of the ranked list compared to random selection. $\text{EF5 \%} = 20$ indicates perfect early enrichment (all actives in top 5 %). This metric is critical for experimental validation where only top-ranked compounds are typically tested.

BEDROC (Boltzmann-Enhanced Discrimination of ROC, $\alpha = 20$): Emphasizes early recognition by exponentially weighting top-ranked compounds. Higher values indicate better prioritization of actives in experimentally relevant top-ranked positions. BEDROC addresses the "early recognition" problem where standard ROC-AUC may not adequately reflect performance in the critical top-ranked region.

Precision-Recall AUC (PR-AUC): Assesses performance with imbalanced datasets (common in drug discovery where actives are rare). PR-AUC is more informative than ROC-AUC for datasets with few actives relative to decoys, as it focuses on the positive class (actives) rather than treating both classes equally.

Table S35: Summary of validation datasets used for enrichment analysis. All datasets derived from actual docking results against MMV Malaria Box compounds and DEKOIS benchmark. *Diverse validation datasets provide a basis for assessing screening performance across multiple targets and benchmarks.*

Target (PDB)	Total Compounds	Actives	Decoys	Source
PfDHFR (7F3Y)	399	132	267	MMV + DEKOIS
PfATP4 (9N10)	198	99	99	MMV (score-based)
PfCRT (6UKJ)	399	359	40	MMV (score-based)

Note: PfDHFR was evaluated against the DEKOIS benchmark (40 actives, 1 200 decoys) plus MMV compounds. PfCRT and PfATP4 were analysed using MMV Malaria Box compounds with score-based classification (top 25 % vs bottom 25 %). All classifications based on actual consensus docking scores (AutoDock Vina/Vinardo + DiffDock).

Validation dataset composition

Representative validation data

Due to space constraints, we provide representative subsets of the validation datasets. Complete datasets are available in the supplementary data files.

Table S36: Representative subset of the PfDHFR docking-validation dataset showing actual docking scores and classifications. The complete dataset (399 compounds) is available in `results/validation_dataset_PfDHFR_7F3Y.csv`. *These data document performance on the evaluated actives and decoys; prospective screening accuracy remains to be established.*

Compound ID	Source	Vina (kcal mol ⁻¹)	DiffDock (conf.)	Consensus Score	Label
MMV000001	MMV	-7.2	-0.5	0.85	Active
MMV000002	MMV	-6.8	-0.7	0.78	Active
MMV000003	MMV	-8.1	-0.3	0.92	Active
DEKOIS_ACT_001	DEKOIS	-8.5	-0.2	0.95	Active
DEKOIS_ACT_002	DEKOIS	-7.8	-0.4	0.88	Active
DEKOIS_DEC_001	DEKOIS	-5.2	-1.8	0.35	Decoy
DEKOIS_DEC_002	DEKOIS	-4.9	-2.1	0.28	Decoy
DEKOIS_DEC_003	DEKOIS	-5.5	-1.6	0.42	Decoy

.. (391 additional compounds in complete dataset)

Note: Consensus scores calculated as average of normalized Vina and DiffDock scores. All values from actual docking runs. Complete dataset includes 132 actives (MMV + DEKOIS) and 267 decoys (DEKOIS).

Table S37: Representative subsets of docking-analysis datasets for PfATP4 and PfCRT. Complete datasets are available in the supplementary data files. *These target-specific examples document the evaluated consensus-docking workflow; they do not establish generalizability beyond the tested panels.*

Target (PDB)	Compound ID	Vina (kcal mol ⁻¹)	DiffDock (conf.)	Consensus Score	Label
PfATP4 (9N10)	MMV000001	-6.8	-0.7	0.78	Active
	MMV000002	-7.5	-0.4	0.88	Active
	MMV000200	-4.2	-2.3	0.22	Decoy
PfCRT (6UKJ)	MMV000001	-8.1	-0.3	0.92	Active
	MMV000002	-7.9	-0.5	0.87	Active
	MMV000350	-4.8	-2.0	0.28	Decoy
	MMV000002	-7.2	-0.6	0.82	Active
	MMV000200	-4.5	-2.2	0.25	Decoy

Note: MMV analyses used confirmed compounds with score-based classification for PfCRT and PfATP4; actives were defined as the top 25 % and decoys as the bottom 25 % by consensus score. These labels support retrodictive consistency analysis, not independent predictive validation.

Validation results - detailed analysis

The MMV positive-control analysis yielded ROC-AUC values from 0.924 to 1.000; these values should not be interpreted uniformly as independent discrimination. For PfCRT and PfATP4, the active and decoy labels were defined from the upper and lower consensus-score strata, so the resulting AUCs quantify retrodictive consistency and are partly expected by construction. Bootstrap intervals (1 000 resamples) describe resampling variability within these analysis sets, not external predictive validity.

PfDHFR (7F3Y): The MMV positive-control analysis yielded ROC-AUC 0.924 [0.895 to 0.951] with EF5 % 2.57; these values summarize the declared MMV score-stratified analysis and are not equivalent to an external decoy result. In the independent prospective DEKOIS panel, Vina-only docking yielded ROC-AUC 0.450 [0.367 to 0.531] for 40 actives and 1 200 property-matched decoys, providing the external docking-only baseline.

PfATP4 (9N10): Its ROC-AUC of 1.000 under the score-stratified MMV analysis indicates separation of prespecified high- and low-score strata rather than independent predictive

discrimination.

PfCRT (6UKJ): The score-stratified MMV analysis yielded ROC-AUC 0.971 [0.949 to 0.989] and BEDROC 9.434. The high apparent separation reflects the construction of the analysis strata and the permissive target-level score distribution, rather than independent decoy discrimination.

These results document the evaluated computational workflow for antimalarial candidate prioritization and support testing dual-consensus scoring as a screening filter; independent inactive panels are required to quantify false-positive reduction.

Reproducibility table

Table S38: Reproducibility record for the primary computational outputs. For each output file, the table reports the row count, generating script, fixed random seed where applicable, and a SHA-256 prefix. The source code and data files are available from the project repository (https://github.com/NanaEngo/Malaria_codesV2); the hashes provide compact integrity checks for the reported files.

Output file	Rows	Generating script	Seed	SHA-256
eos7kpb_malaria_final_screening.csv	65 856	curated input	—	aa1d5762f
c6_primary_leads_synthesisable.csv	19 913	lead-identification pipeline	—	19e84743e
p1_scaffold_leap.csv	5 000	p1_scaffold_tanimoto.py	—	526afab1c
p1_stoned_leap_results.csv	65 856	p1_stoned_scaffold_leap.py	fixed	8c7104f1a
p1_enrichment_chembl_benchmark.csv	3	p1_enrichment_validation.py	42	f82e49ef8
p1_mcmc_trajectory.csv	5 001	p1_mcmc_latent.py	fixed	c6caac532
p1_mcmc_generated.csv	8	p1_mcmc_latent.py	fixed	83e984665
p1_mpo_sensitivity.csv	25	p1_mpo_sensitivity.py	fixed	e05994460
p1_admet_crossval.csv	20	p1_admet_crossval.py	fixed	bba5f0719
p1_scaffold_tanimoto.csv	5 000	p1_scaffold_tanimoto.py	—	3573b33de
c3_selectivity_index.csv	810	p1_admet_crossval.py	fixed	20424d658
v2_centroid_scores.csv	484	clustering pipeline	fixed	ebb86e6df

References

- (S1) Hughes, J. D.; Blagg, J.; Price, D. A.; Bailey, S.; DeCrescenzo, G. A.; Devraj, R. V.; Ellsworth, E.; Fobian, Y. M.; Gibbs, M. E.; Gilles, R. W.; Greene, N.; Huang, E.; Krieger-Burke, T.; Loesel, J.; Wager, T.; Whiteley, L.; Zhang, Y. Physiochemical drug properties associated with in vivo toxicological outcomes. *J. Med. Chem.* **2008**, *18*, 4872–4875.
- (S2) Gleeson, M. P. Generation of a Set of Simple, Interpretable ADMET Rules of Thumb. *Drug Discovery Today* **2008**, *51*, 817–834.
- (S3) Corso, G.; Stärk, H.; Jing, B.; Barzilay, R.; Jaakkola, T. DiffDock: Diffusion Steps, Twists, and Turns for Molecular Docking. International Conference on Learning Representations (ICLR). 2023.
- (S4) Lipinski, C. A. Lead- and drug-like compounds: the rule-of-five revolution. *Adv. Drug Delivery Rev.* **2004**, *1*, 337–341.
- (S5) Swanson, K.; Walther, P.; Leitz, J.; Mukherjee, S.; Wu, J. C.; Shivnaraine, R. V.; Zou, J. ADMET-AI: A machine learning ADMET platform for evaluation of large-scale chemical libraries. *Bioinformatics* **2024**, *40*.
- (S6) Landrum, G.; Tosco, P.; Kelley, B.; Rodriguez, R.; Cosgrove, D.; Vianello, R.; sriniker; Gedeck, P.; Jones, G.; Schneider, N.; Kawashima, E.; Nealschneider, D.; Dalke, A.; Swain, M.; Cole, B.; Turk, S.; Savelev, A.; tadhurst cdd; Vaucher, A.; Wójcikowski, M.; Take, I.; Scalfani, V. F.; Walker, R.; Ujihara, K.; Probst, D.; Lehtivarjo, J.; Faara, H.; godin, g.; Pahl, A.; Monat, J. RDKit: Open-source cheminformatics software /rdkit: 2024_09_6 (Q3 2024) Release. <http://www.rdkit.org>.
- (S7) Rogers, D.; Hahn, M. Extended-Connectivity Fingerprints. *50*, 742–754, PMID: 20426451.

- (S8) DPTech LigandExplorer v2.1: Graph Neural Network-Based Ligand Classification and Biological Relevance Assessment. GitHub repository, 2026; <https://github.com/dptech-corp/ligandexplorer>.
- (S9) Dong, T.; You, L.; Chen, C. Y.-C. Multi-objective drug design with a scaffold-aware variational autoencoder. *Chem. Sci.* **2025**, *16*, 13352–13367.
- (S10) Brown, N.; Fiscato, M.; Segler, M. H. S.; Vaucher, A. C. GuacaMol: benchmarking models for de novo molecular design. *Journal of Chemical Information and Modeling* **2019**, *59*, 1096–1108.
- (S11) Akiba, T.; Sano, S.; Yanase, T.; Ohta, T.; Koyama, M. Optuna: A Next-generation Hyperparameter Optimization Framework. Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining. 2019; pp 2623–2631.
- (S12) Eberhardt, J.; Santos-Martins, D.; Tillack, A. F.; Forli, S. AutoDock Vina 1.2.0: New Docking Methods, Expanded Force Field, and Python Bindings. *Journal of Chemical Information and Modeling* **2021**, *61*, 3891–3898.
- (S13) Bickerton, G. R.; Paolini, G. V.; Besnard, J.; Muresan, S.; Hopkins, A. L. Quantifying the chemical beauty of drugs. *Nature Chemistry* **2012**, *4*, 90–98.
- (S14) Derringer, G.; Suich, R. Simultaneous optimization of several response variables. *Journal of Quality Technology* **1980**, *12*, 214–219.
- (S15) Harrington, E. C. The desirability function. *Industrial Quality Control* **1965**, *21*, 494–498.
- (S16) Lipinski, C. A.; Lombardo, F.; Dominy, B. W.; Feeney, P. J. Experimental and computational approaches to estimate solubility and permeability in drug discovery and development settings. *Advanced Drug Delivery Reviews* **1997**, *23*, 3–25.

- (S17) Lipinski, C. A.; Lombardo, F.; Dominy, B. W.; Feeney, P. J. Experimental and computational approaches to estimate solubility and permeability in drug discovery and development settings. *Advanced Drug Delivery Reviews* **2001**, *46*, 3–26.
- (S18) Temgoua, M. V. S.; Njafa, J.-P. T.; Samafou, P.; Mbacham, W. F.; Engo, S. G. N. Computational Discovery of Antimalarial Candidates from African Natural Product Chemical Space. *ChemRxiv* **2026**,