

Supporting Information

Myke Vital Sao Temgoua Jean-Pierre Tchapet Njafa Serge Guy Nana Engo
Penabei Samafou Wilfred Fon Mbacham

1 Orthogonal LISH mechanism-of-action benchmark

The main text summarises a phenotype-only benchmark on the public LISH mechanism-of-action (MoA) release; the protocol and results are detailed here. The benchmark addresses pharmacology rather than chemistry: it evaluates how well gene-expression and viability readouts predict a drug’s MoA, and it says nothing about whether the molecular representations examined in the main analysis capture that pharmacology, because no structural information links the two data sets.

1.1 Data and preparation

We drew on the public LISH MoA training release [1], available from the challenge repository at https://github.com/LISHarvard/moa_challenge (accessed 12 August 2026), and treated it as an independent, phenotype-only data set. After aggregating repeated assay observations to the drug level, the working table retained 772 gene-expression features, 100 cell-viability features, and three experimental-condition summaries (exposure time, high-dose fraction, treatment-type fraction).

The scored endpoint comprises 206 MoA labels across 3 289 drugs [2], each label assigned at the drug level from the observed MoA-associated annotations. The public release carries no versioned mapping from drug identifier to SMILES, so the data set could not be used to train or evaluate the ECFP4 fingerprint, the graph neural networks, or the transformer; the benchmark stands on its own as a phenotype-only reference. A prespecified sensitivity check removed vehicle-control samples, leaving 3 288 drugs, and every fold record was checked for missing or non-finite metrics before the results were interpreted. Because MoA annotation reflects observed association rather than a measured physical interaction, the endpoint is referred to throughout as MoA-associated prediction rather than as target engagement [2].

1.2 Protocol

For each of the 206 labels an unweighted logistic model was fitted independently and evaluated over five random seeds and five drug-grouped folds. Mean column-wise log loss served as the primary metric; macro-AUPRC [3], macro-AUROC [4], mean Brier score and expected calibration error [5] accompany it as secondary summaries of the same fold grid.

1.3 Results

Table 1 reports the phenotype-only baseline, averaged over the 25 fold–seed evaluations, with vehicle controls retained. Exclusion of those controls left macro-AUROC essentially unchanged (−0.0018) and moved every other metric by less than 0.002 in absolute terms, so the baseline’s behaviour does not depend on whether vehicle-treated samples remain in the training data.

Table 1: Phenotype-only LISH baseline metrics, mean across 25 fold–seed evaluations (five seeds $\times$ five drug-grouped folds).

Metric	With controls	Without controls	Δ
Mean column-wise log loss	0.023 78	0.023 89	+0.000 11
Macro-AUROC	0.643 5	0.641 7	−0.001 77
Macro-AUPRC	0.142 8	0.141 2	−0.001 52
Mean Brier score	0.003 74	0.003 75	+0.000 01
Expected calibration error	0.003 76	0.003 81	+0.000 05

1.4 Scope of the comparison

These metrics characterise the phenotype-only reference on its own terms; they cannot be set numerically against the molecular ROC-AUC results in the main text, because the two analyses differ in task, labels, features, unit of aggregation and primary metric. The labels themselves record MoA association rather than a demonstrated physical interaction with a target. The LISH result is therefore best read as an orthogonal pharmacological reference, together with a check on its sensitivity to vehicle controls, rather than as evidence bearing on the molecular models or on causal mechanism.

2 Post-hoc calibration metrics for the molecular benchmark

Calibration metrics were computed after the fact from the archived fold-level test predictions of the extended campaign, pooling five seeds and five folds per configuration into 99 180 test records each; no model was retrained for this analysis. The resulting expected calibration error (ECE) and maximum calibration error (MCE), computed with 10 equal-width bins, and the Brier score (mean squared error between predicted probability and label) describe the calibration of the deployed checkpoints as a group, not of any single fold or seed in isolation, because the predictions come from checkpoints already fitted during the extended campaign and are pooled across seeds and folds.

Table 2: Post-hoc calibration metrics (pooled over 25 fold–seed evaluations per configuration). ECE and MCE rise from the random split to the scaffold and novel-partition splits for every model family, indicating that confidence degrades under scaffold shift even when ranking (ROC-AUC) is only moderately affected.

Split	Model	ECE	MCE	Brier
Canonical random	GIN	0.029 0	0.069 9	0.102 55
	GIN-TFP (native)	0.034 4	0.071 6	0.105 24
	GIN-TNE (native)	0.036 7	0.093 6	0.112 86
	ChemBERTa	0.057 8	0.144 0	0.104 66
Canonical scaffold	GIN	0.065 4	0.162 3	0.148 59
	GIN-TFP (native)	0.074 1	0.175 4	0.147 69
	GIN-TNE (native)	0.080 7	0.194 6	0.151 08
	ChemBERTa	0.122 5	0.254 6	0.169 09
Novel partitions (range)	GIN	0.0777–0.0888	0.1814–0.1854	0.15359–0.15919
	GIN-TFP (native)	0.0741–0.0895	0.1737–0.1890	0.14922–0.15564
	GIN-TNE (native)	0.0798–0.0820	0.1813–0.2078	0.15305–0.15539
	ChemBERTa	0.1172–0.1320	0.2401–0.2715	0.16699–0.17508

The pattern is consistent across model families: ECE roughly doubles from the random split

to the scaffold split and widens further on the novel partitions. ChemBERTa, already the arm with the widest fold-level spread in ROC-AUC, also carries the largest calibration error under scaffold shift. These pooled estimates refine the ranking analysis without revising it; they leave the canonical ROC-AUC results untouched and should not be read as calibrated probabilities ready for prospective use. Per-configuration reliability diagrams are available in the released data package.

2.1 Calibration and ranking on the high-RRS polypharmacology subset

Calibration and ROC-AUC analyses were also stratified by resistance-relevant score (RRS), the docking-based classification from the companion P2 Set-C study that uses docking-RRS $\geq 80\%$ on available targets. Applying that classification [6], six candidates with an A*/A RRS class on the targets for which docking scores were available (PP-01, PP-02, PP-06, PP-11, PP-13, PP-15) were identified, together contributing 30 test samples per configuration out of the 99 180 pooled records. This available-target definition is broader than a stricter, complete two-target definition (PP-01 and PP-15 only, $n = 2$), because it also admits candidates that lack a score on one target but retain a high RRS on the other. Calibration broken down by RRS class, under both the complete and available-target definitions, is provided in the released data package.

Table 3: Available-target polypharmacology subset ROC-AUC on the scaffold split (six A*/A-class candidates, 30 test samples per configuration against 99 150 single-target samples). Δ AUC is polypharmacology minus single-target. The small candidate count precludes a bootstrap 95% interval, so values are reported descriptively.

Model	Polypharma AUC	Single-target AUC	Δ AUC
ECFP4-RF	0.616	0.789	−0.173
GIN	0.144	0.812	−0.668
GIN-TFP (native)	0.296	0.817	−0.521
GIN-TNE (native)	0.224	0.815	−0.591
ChemBERTa	0.616	0.789	−0.173

On the scaffold split, ROC-AUC on the available-target polypharmacology subset falls below the single-target value for every model, with deficits ranging from -0.173 to -0.668 . ChemBERTa and ECFP4-RF hold up best (a deficit of -0.173 each), while GIN, GIN-TFP and GIN-TNE lose considerably more (-0.52 to -0.67). The stricter complete two-target subset, built from only two candidates, is noisier still (AUC ranging 0.00–0.84) and shows even larger deficits for most models. These numbers are descriptive only: the small candidate count prevents any inferential claim, and the honest-negative scaffold result that anchors the main analysis stands unchanged.

3 Light GNN hyperparameter sensitivity on the scaffold split

The canonical negative result could, in principle, reflect a poor choice of GIN capacity rather than a genuine limit of the representation. To probe this possibility the base GIN was retrained on the frozen canonical scaffold split at two settings that bracket the canonical configuration (hidden dimension 128, dropout 0.1): a smaller model (hidden 64, dropout 0.2) and a larger one (hidden 256, dropout 0.1), keeping every other element of the protocol identical—five seeds, five folds, the same optimiser constants, epoch budget and early-stopping rule.

All three configurations remain clearly below the ECFP4-RF [7] scaffold reference (ROC AUC 0.8300), and the spread across them (0.8000–0.8081) is small relative to the gap that separates any of them from the fingerprint baseline (0.022–0.030). Doubling or halving the hidden capacity, or changing the dropout rate, does not close that gap on this panel. The

Table 4: Base-GIN sensitivity to hidden dimension and dropout on the scaffold split (25 fold-seed records per configuration). The canonical configuration (hidden 128, dropout 0.1) reached ROC AUC 0.8047.

Configuration	ROC AUC (mean of 25)	Per-seed SD	Fold-level SD
Hidden 64, dropout 0.2	0.8081	0.0118	0.0386
Hidden 128, dropout 0.1 (canonical)	0.8047	0.0141	0.0395
Hidden 256, dropout 0.1	0.8000	0.0160	0.0370

check bounds the possibility that some untested architecture or training budget would behave differently without eliminating it; taken together with the canonical result, it supports reading the negative finding as a property of the representation comparison itself rather than of one particular capacity choice.

4 Paired differences against the ECFP4-RF reference

The main text reports each model-minus-reference difference alongside a permutation-based interval. Table 5 complements that summary with the paired seed-level contrasts and the 95% intervals derived from the paired t statistics (five per-seed means per arm). All eight contrasts are negative and exclude zero at the 95% level after Benjamini-Hochberg correction [8].

Table 5: Paired model-minus-ECFP4-RF AUC differences with 95% intervals, derived from the paired seed-level statistics.

Split	Model	Δ AUC	95% interval	Adjusted p
Random	GIN	-0.0335	-0.0368—0.0302	9.9×10^{-6}
Random	GIN-TFP	-0.0349	-0.0365—0.0333	6.0×10^{-7}
Random	GIN-TNE	-0.0515	-0.0539—0.0491	6.0×10^{-7}
Random	ChemBERTa	-0.0312	-0.0325—0.0299	6.0×10^{-7}
Scaffold	GIN	-0.0253	-0.0434—0.0072	0.0330
Scaffold	GIN-TFP	-0.0162	-0.0290—0.0034	0.0330
Scaffold	GIN-TNE	-0.0210	-0.0401—0.0019	0.0378
Scaffold	ChemBERTa	-0.0433	-0.0499—0.0367	0.0002

These paired- t -based intervals are wider than the permutation intervals reported for the same four scaffold contrasts in the main text, as expected from the two different construction methods. The two tables agree on point estimates, sign and exclusion of zero at the same four Δ AUC values. The intervals reported here describe the evaluated seed-fold grid and should not be extended to other panels, splits or architectures.

References

- [1] Jianming Ma, Robert P. Sheridan, Andy Liaw, George E. Dahl, and Vladimir Svetnik. Deep neural net chemical activity prediction using SMILES. In *Neural Information Processing Systems Workshop on Machine Learning for Drug Discovery*, 2019.
- [2] Maria-Anna Trapotsi, Layla Hosseini-Gerami, and Andreas Bender. Computational analyses of mechanism of action (moa): data, methods and integration. *RSC Chemical Biology*, 3(2): 170–200, 2022. doi: 10.1039/D1CB00069A.

- [3] Jesse Davis and Mark Goadrich. The relationship between precision-recall and ROC curves. In *Proceedings of the 23rd International Conference on Machine Learning*, pages 233–240, 2006.
- [4] James A. Hanley and Barbara J. McNeil. The meaning and use of the area under a receiver operating characteristic (ROC) curve. *Radiology*, 143:29–36, 1982.
- [5] Mahdi Pakdaman Naeini, Gregory F. Cooper, and Milos Hauskrecht. Obtaining well calibrated probabilities using bayesian binning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 29, 2015.
- [6] Myke Vital Sao Temgoua, Jean-Pierre Tchapet Njafa, Penabei Samafou, and Wilfred Fon Mbacham. Molecular dynamics validation of polypharmacological antimalarial leads targeting resistance mutations. *Journal of Chemical Information and Modeling*, 2026. Companion manuscript (Paper 2), submitted for publication in Journal of Chemical Information and Modeling.
- [7] D. Rogers and M. Hahn. Extended-connectivity fingerprints. *Journal of Chemical Information and Modeling*, 50(5):742–754, 2010. doi: 10.1021/ci100050t.
- [8] Yoav Benjamini and Yosef Hochberg. Controlling the false discovery rate: A practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society: Series B*, 57: 289–300, 1995.